



AI VIET NAM


@aivietnam.edu.vn

NLP Basic – 08

Pre-trained Word Vectors

AI VIET NAM

Nguyễn Quốc Thái

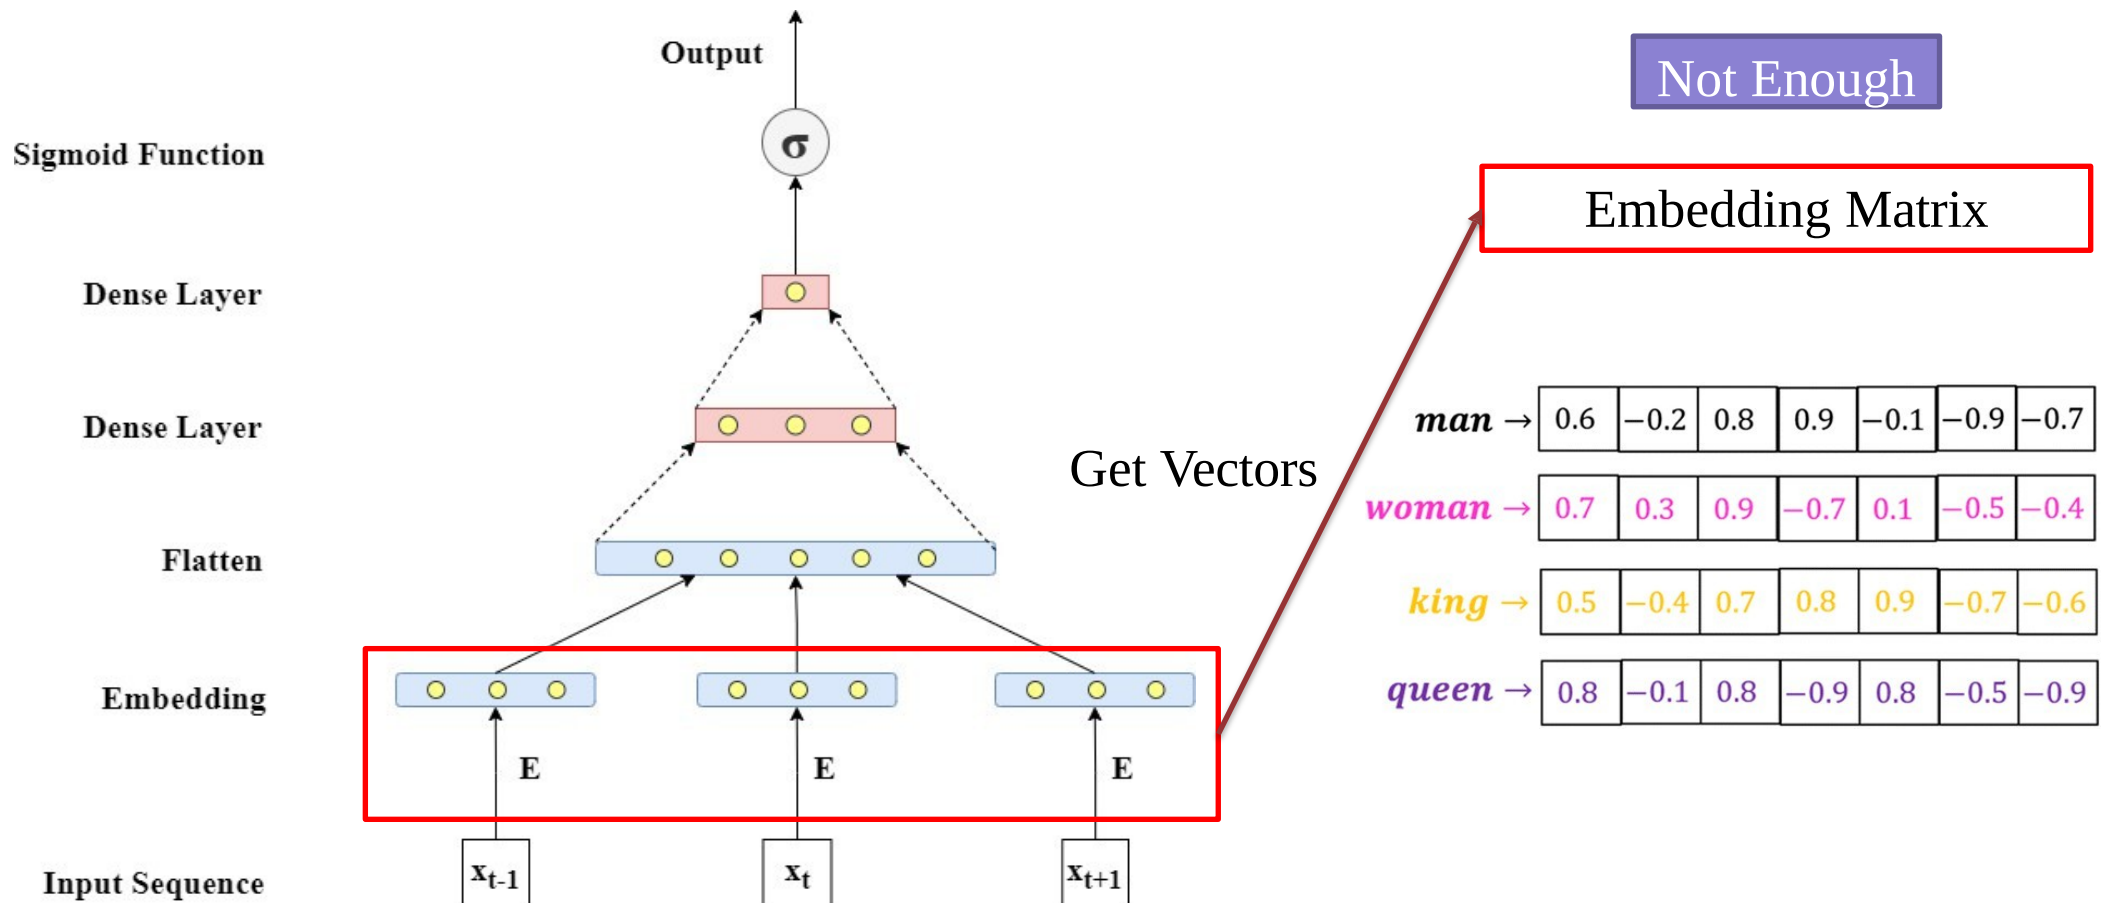


CONTENT

1	Pre-trained Word Vectors
2	Text Classification
3	Summary

1 – Pretrained Word Vectors

❖ Pre-trained Language Models (LMs)



1 – Pretrained Word Vectors

❖ Pre-trained Language Models (LMs)

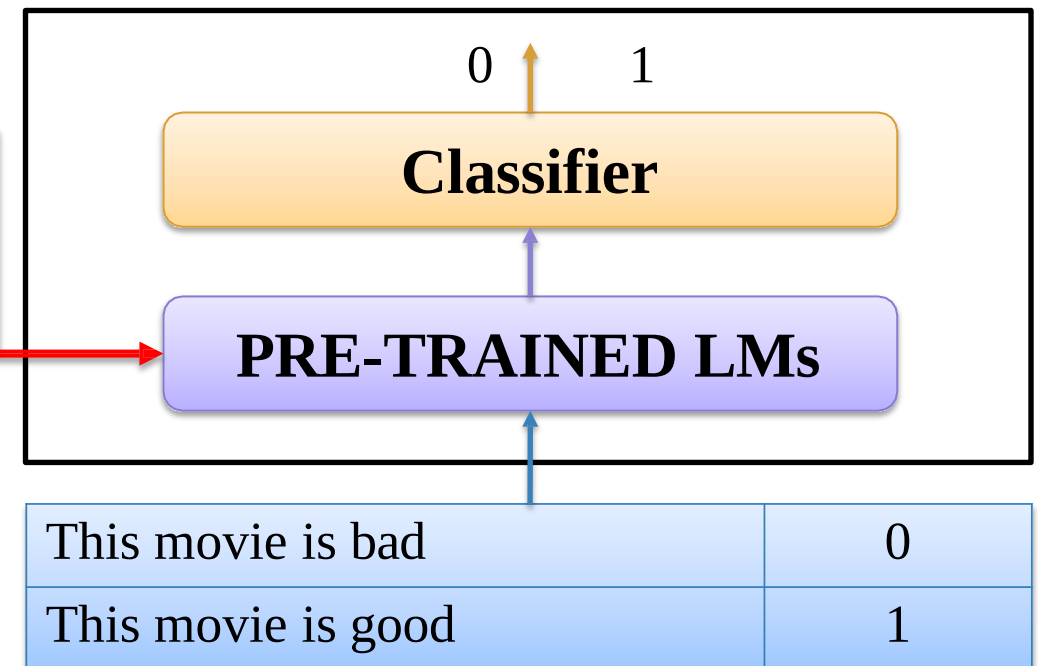
Objective: Language Model

Supervised: Text Classification

Model


LANGUAGE MODEL

Data







1 – Pretrained Word Vectors

❖ Pre-trained Language Models (LMs)



1 – Pretrained Word Vectors

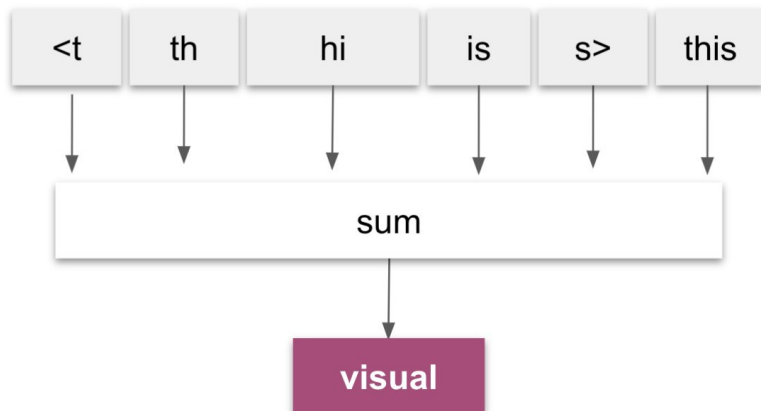
❖ Word2Vec




1 – Pretrained Word Vectors

❖ Fasttext

fastText



Word2Vec

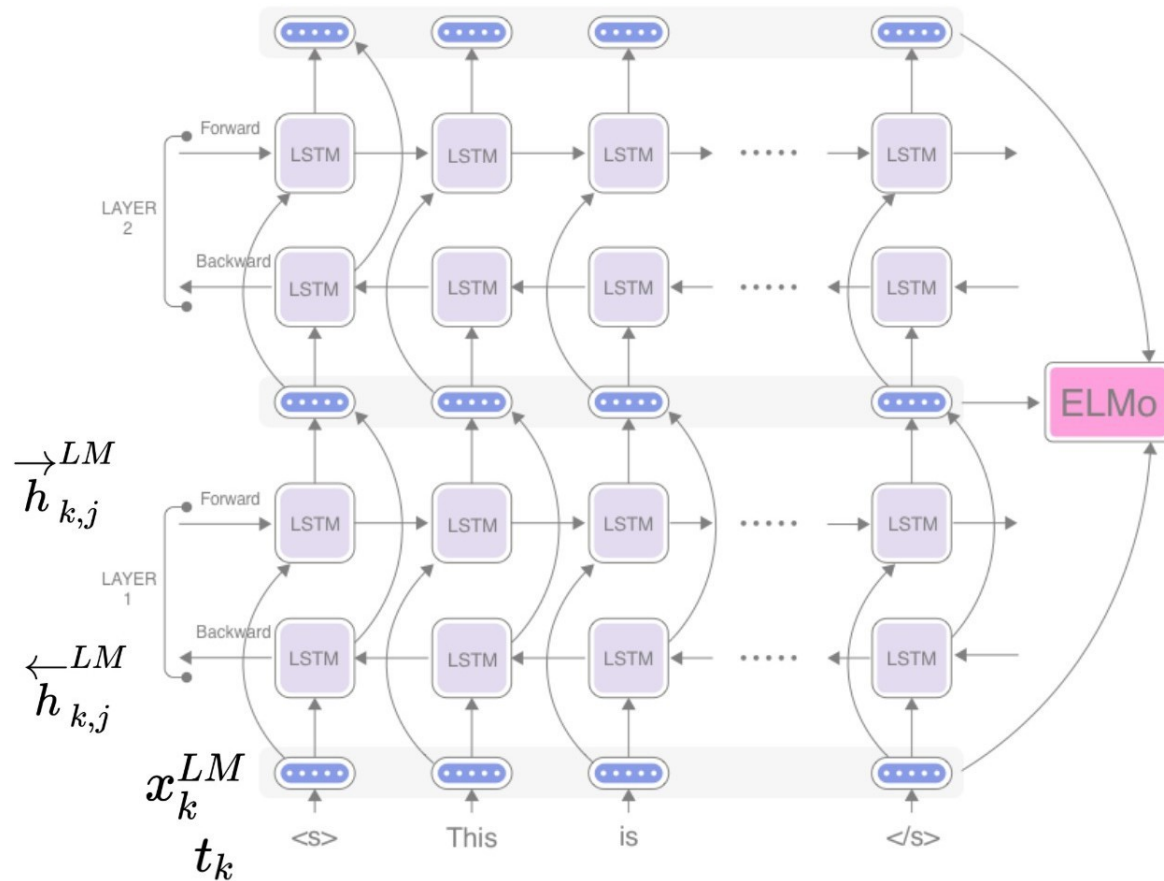


Perplexity

	Cs	DE	ES	FR	RU
Vocab. size	46k	37k	27k	25k	63k
CLBL	465	296	200	225	304
CANLM	371	239	165	184	261
LSTM	366	222	157	173	262
sg	339	216	150	162	237
sisg	312	206	145	159	206

1 – Pretrained Word Vectors

❖ ELMo



[Paper](#)
[Pre-trained](#)

TASK	PREVIOUS SOTA	OUR BASELINE	ELMo + BASELINE
SQuAD	Liu et al. (2017) 84.4	81.1	85.8
SNLI	Chen et al. (2017) 88.6	88.0	88.7 ± 0.17
SRL	He et al. (2017) 81.7	81.4	84.6
Coref	Lee et al. (2017) 67.2	67.2	70.4
NER	Peters et al. (2017) 91.93 ± 0.19	90.15	92.22 ± 0.10
SST-5	McCann et al. (2017) 53.7	51.4	54.7 ± 0.5

1 – Pretrained Word Vectors

❖ Glove

- The conditional probability:

$$Q_{ij} = \frac{\exp(u_j^T v_i)}{\sum_{w \in V} \exp(u_j^T w)}$$

- Using co-occurrence probabilities:

$$J = - \sum_i \sum_{j \in \text{context}(i)} Q_{ij} \log X_{ij}$$

NER Task

Model	Dev	Test	ACE	MUC7
Discrete	91.0	85.4	77.4	73.4
SVD	90.8	85.7	77.3	73.7
SVD-S	91.0	85.5	77.6	74.3
SVD-L	90.5	84.8	73.6	71.5
HPCA	92.6	88.7	81.7	80.7
HSMN	90.5	85.7	78.7	74.7
CW	92.2	87.4	81.7	80.2
CBOW	93.1	88.2	82.2	81.1
GloVe	93.2	88.3	82.9	82.2

1 – Pretrained Word Vectors

❖ Pre-trained Glove Embedding

Download pre-trained word vectors

- Pre-trained word vectors. This data is made available under the [Public Domain Dedication and License](http://www.opendatacommons.org/licenses/pddl/1.0/) v1.0 whose full text can be found at: <http://www.opendatacommons.org/licenses/pddl/1.0/>.
 - [Wikipedia 2014](#) + [Gigaword 5](#) (6B tokens, 400K vocab, uncased, 50d, 100d, 200d, & 300d vectors, 822 MB download): [glove.6B.zip](#)
 - Common Crawl (42B tokens, 1.9M vocab, uncased, 300d vectors, 1.75 GB download): [glove.42B.300d.zip](#)
 - Common Crawl (840B tokens, 2.2M vocab, cased, 300d vectors, 2.03 GB download): [glove.840B.300d.zip](#)
 - Twitter (2B tweets, 27B tokens, 1.2M vocab, uncased, 25d, 50d, 100d, & 200d vectors, 1.42 GB download): [glove.twitter.27B.zip](#)
- Ruby [script](#) for preprocessing Twitter data

1 – Pretrained Word Vectors

❖ Pre-trained Glove Embedding

Version:

6B

400K Vocab

50 D

```
embeddings_dict = {}  
with open("/content/glove.6B.50d.txt", 'r') as f:  
    for line in f:  
        values = line.split()  
        word = values[0]  
        vector = np.asarray(values[1:], "float32")  
        embeddings_dict[word] = vector  
def embedding(word):  
    return embeddings_dict.get(word, embeddings_dict.get('unk'))
```

1 – Pretrained Word Vectors

❖ Pre-trained GloVe Embedding

Version:

6B

400K Vocab

50 D



```
embedding('man')
```



```
array([-0.094386,  0.43007 , -0.17224 , -0.45529 ,  1.6447  ,  0.40335 ,  
       -0.37263 ,  0.25071 , -0.10588 ,  0.10778 , -0.10848 ,  0.15181 ,  
       -0.65396 ,  0.55054 ,  0.59591 , -0.46278 ,  0.11847 ,  0.64448 ,  
       -0.70948 ,  0.23947 , -0.82905 ,  1.272   ,  0.033021,  0.2935  ,  
        0.3911  , -2.8094  , -0.70745 ,  0.4106  ,  0.3894  , -0.2913  ,  
        2.6124  , -0.34576 , -0.16832 ,  0.25154 ,  0.31216 ,  0.31639 ,  
        0.12539 , -0.012646,  0.22297 , -0.56585 , -0.086264,  0.62549 ,  
       -0.0576  ,  0.29375 ,  0.66005 , -0.53115 , -0.48233 , -0.97925 ,  
        0.53135 , -0.11725 ], dtype=float32)
```

1 – Pretrained Word Vectors


❖ Pre-trained GloVe Embedding

Version:

6B

400K Vocab

50 D




1 – Pretrained Word Vectors

❖ Pre-trained GloVe Embedding

Find Synonyms

Find Analogies

```
def euclidean_distance(x, y):  
    return np.sqrt(np.sum((x - y) ** 2))  
  
def cosine_similarity(x, y):  
    return np.dot(x, y) / (np.sqrt(np.dot(x, x)) * np.sqrt(np.dot(y, y)))
```



1 – Pretrained Word Vectors

❖ Pre-trained Glove Embedding

Find Synonyms

Given a word

Find top k synonym words

```
def find_closest_embeddings(target_embedding, top_word, score="cosine"):
    score_dict = {}
    if score == "euclidean":
        for word in embeddings_dict.keys():
            score_dict[word] = euclidean_distance(embedding(word), target_embedding)
        score_dict = sorted(score_dict.items(), key=lambda kv: kv[1])
    if score == "cosine":
        for word in embeddings_dict.keys():
            score_dict[word] = cosine_similarity(embedding(word), target_embedding)
        score_dict = sorted(score_dict.items(), key=lambda kv: kv[1], reverse=True)

    return score_dict[:top_word]
```


1 – Pretrained Word Vectors

❖ Pre-trained Glove Embedding

Find Synonyms

Given a word

Find top k synonym words

```
#Use euclidean distance  
find_closest_embeddings(embedding("man"), 5, score="euclidean")
```

```
[('man', 0.0),  
 ('woman', 2.6026237),  
 ('another', 2.8089325),  
 ('boy', 2.8093922),  
 ('one', 2.9732437)]
```

```
#Use cosine similarity  
find_closest_embeddings(embedding("man"), 5)
```

```
[('man', 1.0),  
 ('woman', 0.8860338),  
 ('boy', 0.8564431),  
 ('another', 0.84528404),  
 ('old', 0.8372183)]
```


1 – Pretrained Word Vectors

❖ Pre-trained GloVe Embedding

Find Analogies

Given 3 words


Find a word with analogies relationship

Example:

“man” : “woman” :: “son” : “daughter”

$a : b :: c : d$

$\text{Vec}(a) + \text{Vec}(d) = \text{Vec}(b) + \text{Vec}(c)$



1 – Pretrained Word Vectors

❖ Pre-trained Glove Embedding

Find Analogies

Given 3 words

Find a word with analogies relationship

Example:

“man” : “woman” :: “son” : “daughter”

$a : b :: c : d$

$\text{Vec}(a) + \text{Vec}(d) = \text{Vec}(b) + \text{Vec}(c)$


```
def find_analogy(word_a, word_b, word_c, score="cosine"):
    x = embedding(word_c) + embedding(word_b) - embedding(word_a)
    return find_closest_embeddings(x, 5, score=score)
```

```
find_analogy("man", "woman", "son", score="euclidean")
```

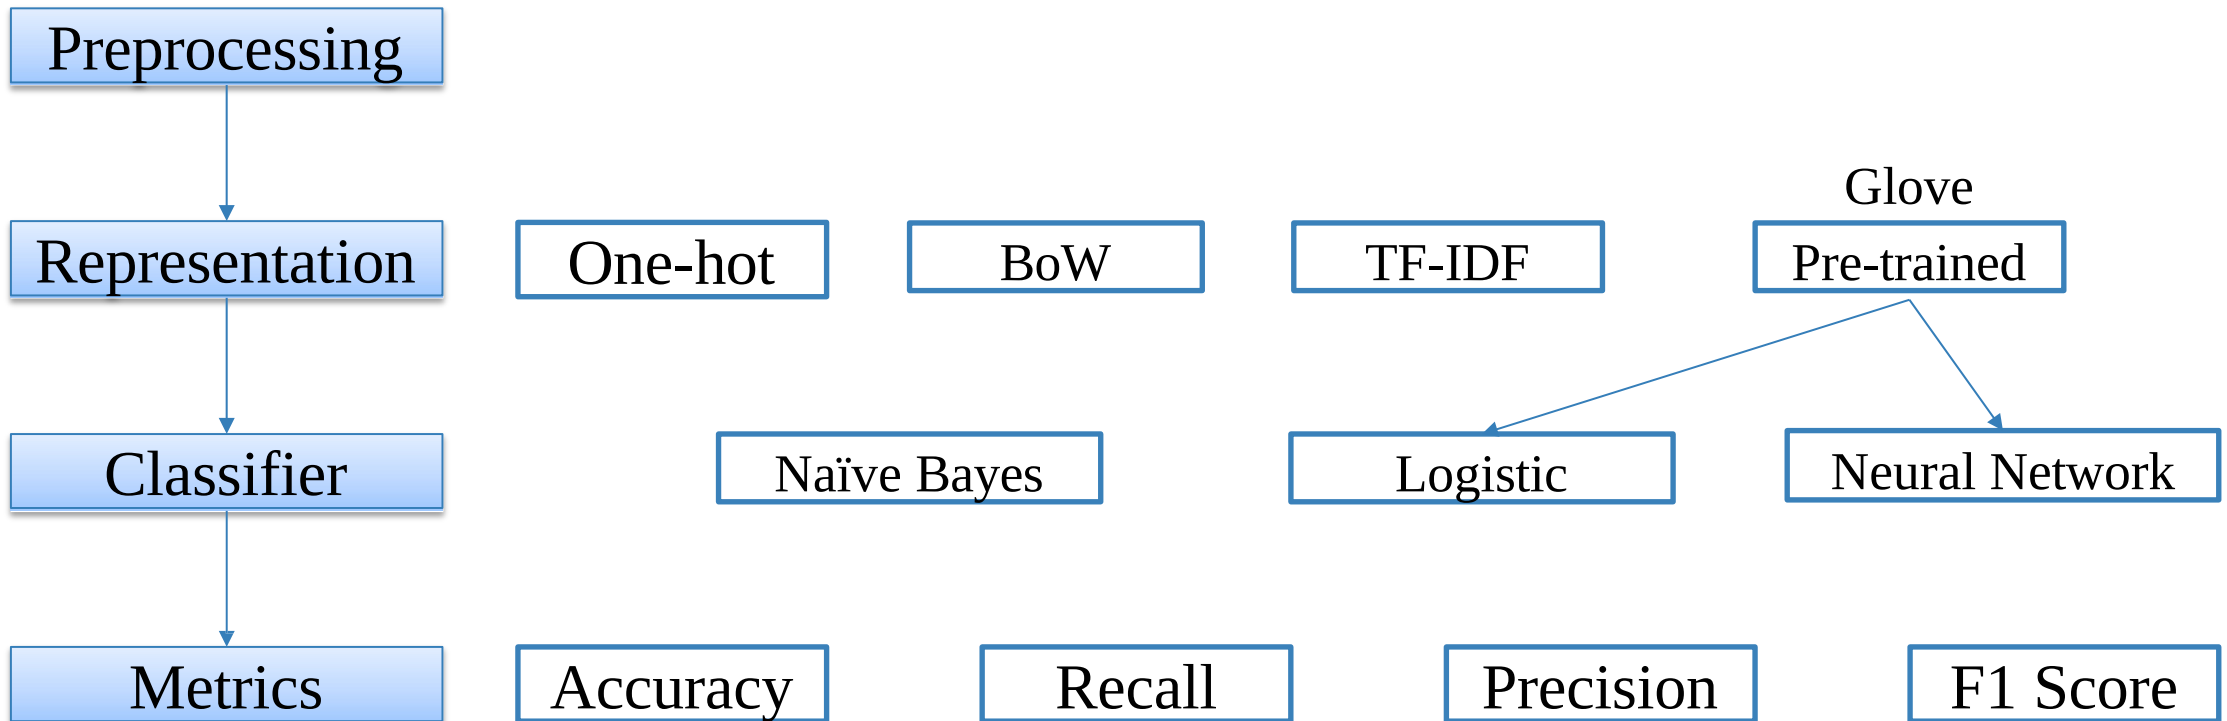
```
[('daughter', 1.5812494),
 ('mother', 2.448467),
 ('wife', 2.4532523),
 ('son', 2.6026237),
 ('father', 2.7388973)]
```

```
find_analogy("man", "woman", "son")
```

```
[('daughter', 0.9658342),
 ('mother', 0.91536117),
 ('wife', 0.9149919),
 ('son', 0.903869),
 ('niece', 0.8937417)]
```

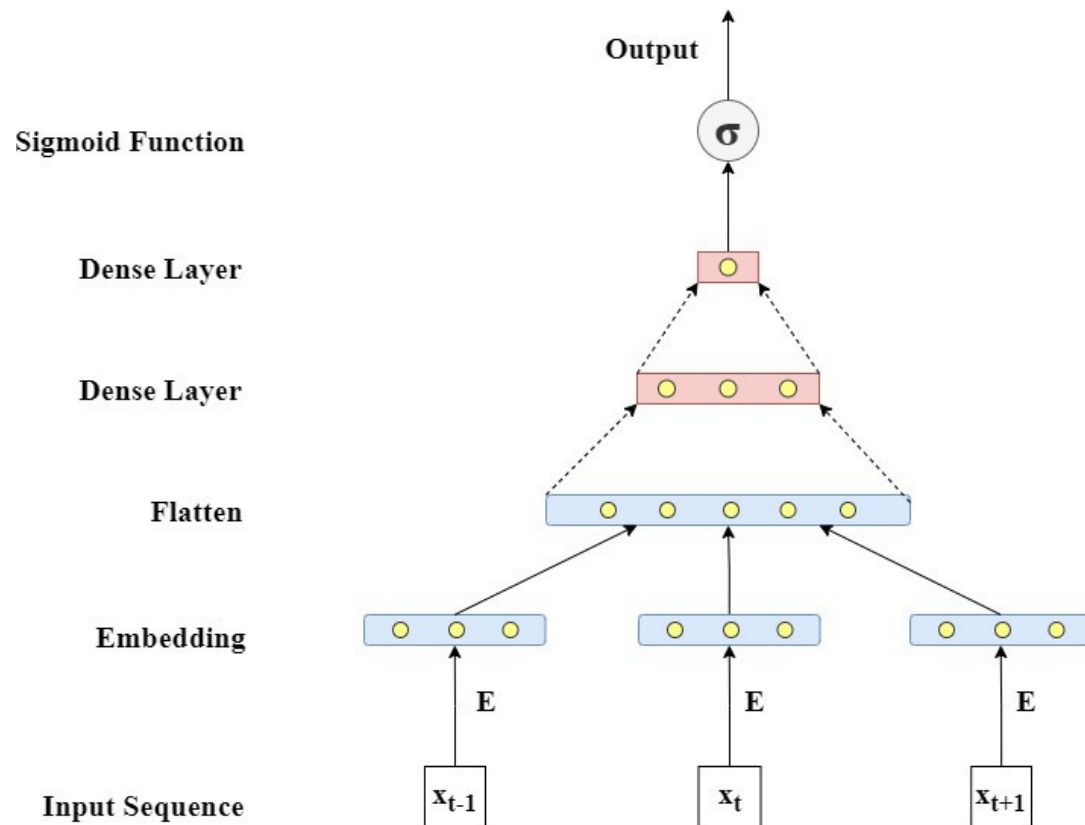


2 – Text Classification



2 – Text Classification

❖ Neural Network



```

1 model_nn = Sequential()
2 model_nn.add(Embedding(vocab_size, embedding_dim, input_length=max_length))
3 model_nn.add(Flatten())
4 model_nn.add(Dense(10, activation='relu'))
5 model_nn.add(Dense(1, activation='sigmoid'))
6 model_nn.compile(
7     loss='binary_crossentropy',
8     optimizer='adam',
9     metrics=['accuracy']
10 )
11 model_nn.summary()

```

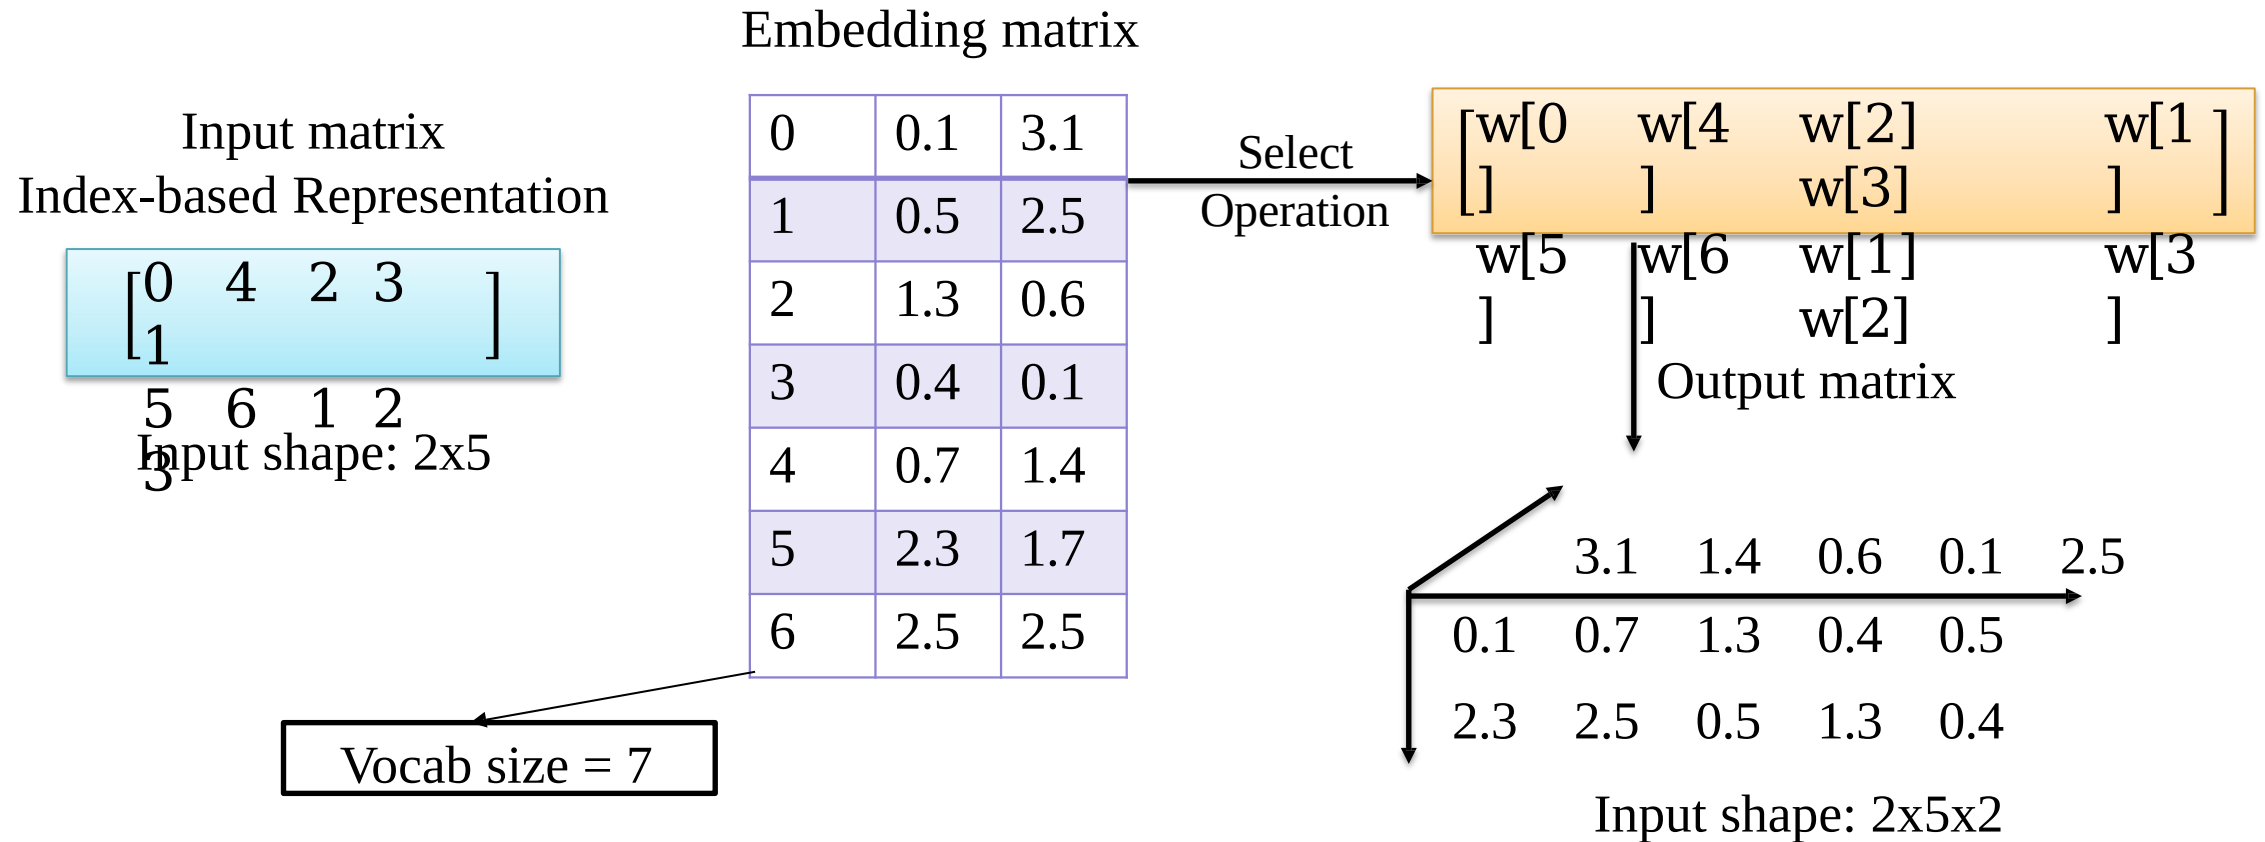
Model: "sequential_2"

Layer (type)	Output Shape	Param #
embedding_2 (Embedding)	(None, 100, 200)	3000000
flatten_2 (Flatten)	(None, 20000)	0
dense_2 (Dense)	(None, 10)	200010
dense_3 (Dense)	(None, 1)	11

Total params: 3,200,021
 Trainable params: 3,200,021
 Non-trainable params: 0


2 – Text Classification

❖ Review: Embedding Layer



2 – Text Classification

❖ Pre-trained Glove Embedding



Embedding matrix

0	<oov>	0.1	3.1
1	<pad>	0.5	2.5
2	<unk>	1.3	0.6
3	neural	0.4	0.1
4	language	0.7	1.4

Glove embedding matrix

0	<oov>	0.1	0.1
1	<pad>	0.5	0.5
2	<unk>	0.3	0.6
3	language	0.7	0.7
4	mưa	0.7	0.4



Final embedding matrix


0	<oov>	0.1	0.1
1	<pad>	0.5	0.5
2	<unk>	0.3	0.6
3	neural	0.0	0.0
4	language	0.7	0.7

2 – Text Classification

❖ Pre-trained Glove Embedding

Glove embedding matrix

0	<oov>	0.1	0.1
1	<pad>	0.5	0.5
2	<unk>	0.3	0.6
3	language	0.7	0.7
4	mưa	0.7	0.4



```
count = 0
embedding_matrix = np.zeros((vocab_size, embedding_dim))

for word, i in tokenizer.word_index.items():
    if i < vocab_size:
        embedding_vector = embeddings_dict.get(word)


        # Words not found in the mebedding index will all l
        if embedding_vector is not None:
            embedding_matrix[i] = embedding_vector
```

2 – Text Classification

❖ Pre-trained Glove Embedding

Final embedding matrix

0	<oov>	0.1	0.1
1	<pad>	0.5	0.5
2	<unk>	0.3	0.6
3	neural	0.0	0.0
4	language	0.7	0.7




```
1 vocab_size = 15000
2 max_length = 100
3 embedding_dim = 200
```

```
1 embeddings_dict = {}
2 with open("/content/glove.6B.200d.txt", 'r') as f:
3     for line in f:
4         values = line.split()
5         word = values[0]
6         vector = np.asarray(values[1:], "float32")
7         embeddings_dict[word] = vector
8 def embedding(word):
9     return embeddings_dict.get(word, embeddings_dict.get('unk'))
```


2 – Text Classification

❖ Pre-trained Glove Embedding



Final embedding matrix

0	<oov>	0.1	0.1
1	<pad>	0.5	0.5
2	<unk>	0.3	0.6
3	neural	0.0	0.0
4	language	0.7	0.7

Update weight

```
model_nn_pre.layers[0].set_weights([embedding_matrix])  
model_nn_pre.layers[0].trainable = False
```

```
model_nn_pre.layers[0].set_weights([embedding_matrix])  
model_nn_pre.layers[0].trainable = True
```

Update weight

3 - Summary

❖ Basic NLP Course

01

Introduction

02

Preprocessing

03

Language Modeling

04

Part Of Speech (POS)

05

Constituency Parsing

06

Basic Vectorization

07

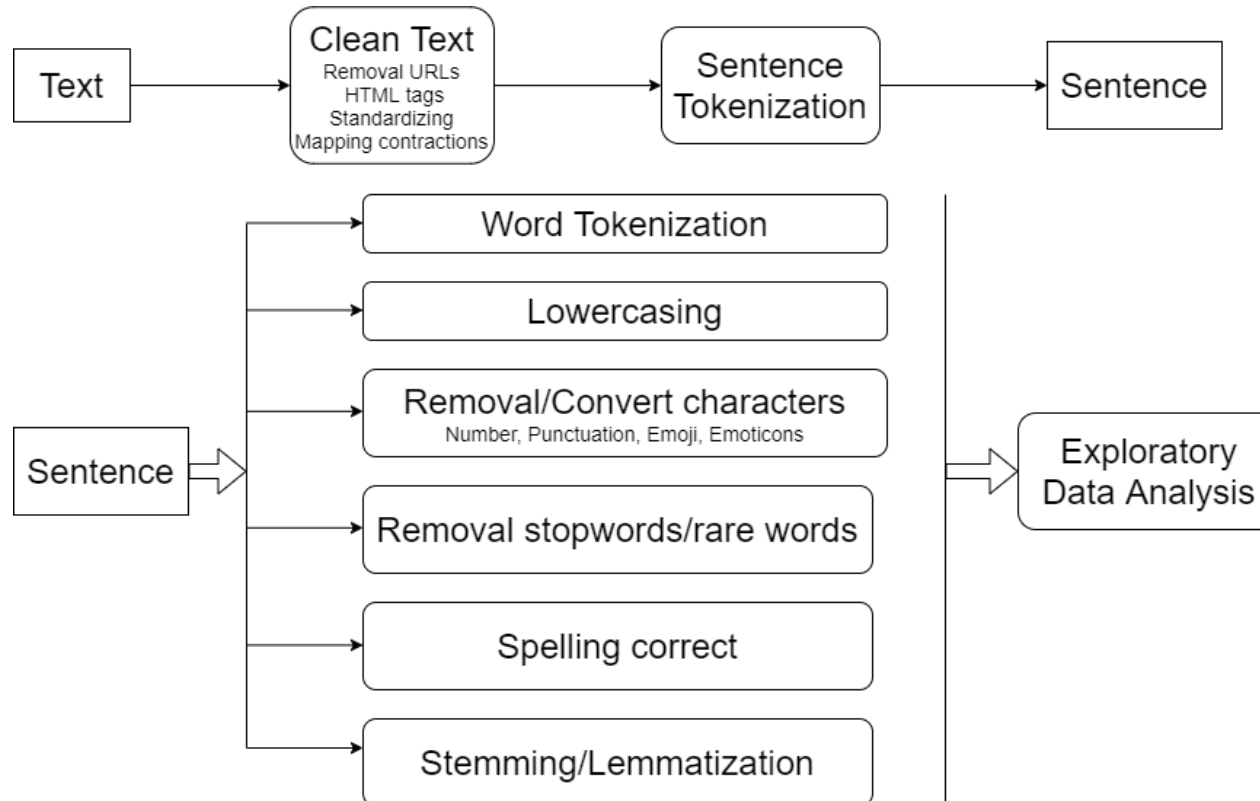
Word2Vec


08

Pretrained Model

3 - Summary

❖ 2 - Preprocessing





3 - Summary


❖ 2 - Preprocessing

Balanced Data




■ Positive ■ Negative

Imbalanced Data




■ Positive ■ Negative



3 - Summary

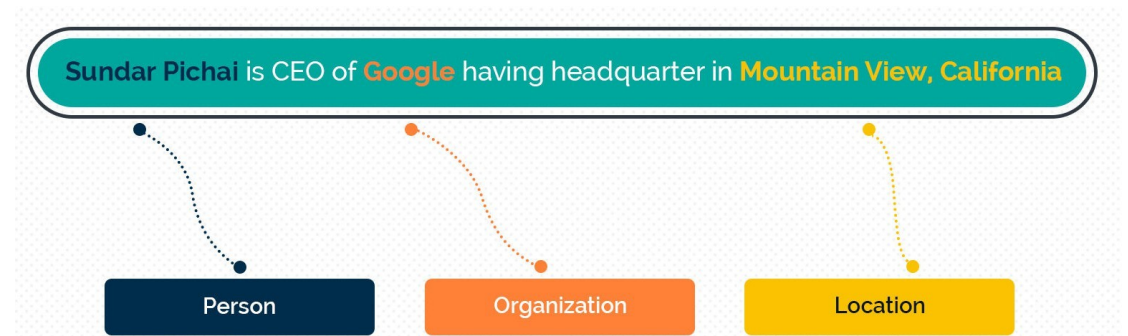

❖ 3 – Language Model



3 - Summary

❖ 4 – POS Tagging - NER

Model: Hidden Markov Model (HMM)



3 - Summary

❖ 5 – Constituency Parsing (CFG)

$G = (T, N, P, S, R)$

T: a set of terminal symbols

N: a set of non-terminal symbols

$P(P, N)$: a set of pre-terminal symbols

S: a start symbol

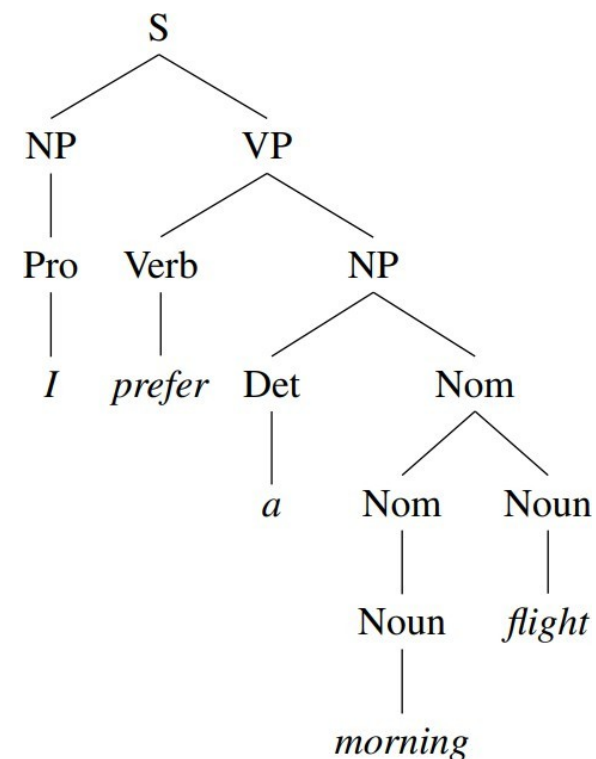
R: a set of rules or productions

$R = \{ \mid N, (T, N) \}$

Grammar in CNF

1. Start S
2. S NP VP
3. NP Det Noun
4. NP NN PP
5. PP Prep NP
6. VP V NP
7. a. VP V Args
b. Args NP PP
8. V ate
9. NP John
10. NP ice-cream, snow
11. Noun ice-cream, pizza
12. Noun table, guy, campus
13. Det the
14. Prep on

CKY



3 - Summary

❖ 5 – Dependency Parsing

A graph $G = (V, A)$

V vertices

$\{w_0=\text{root}, w_1, \dots, w_n\}$

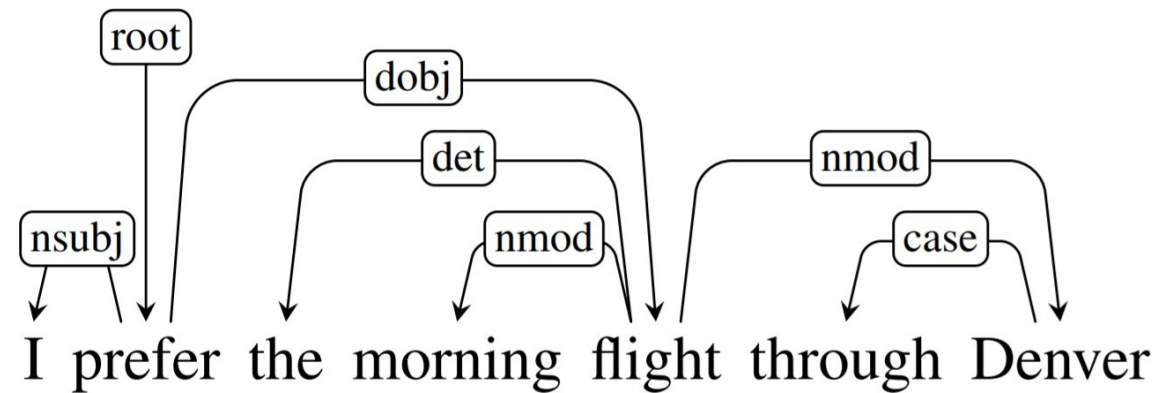
usually one per word in sentence

A arcs

$\{(w_i, r, w_j) : w_i \neq w_j \mid w_i \in V,$

$w_j \in V - w_0, r \in R_x\}$

R_x : a set of all possible dependency relations in x



3 - Summary

❖ 5 – Dependency Parsing

Dependency Tree

a ROOT

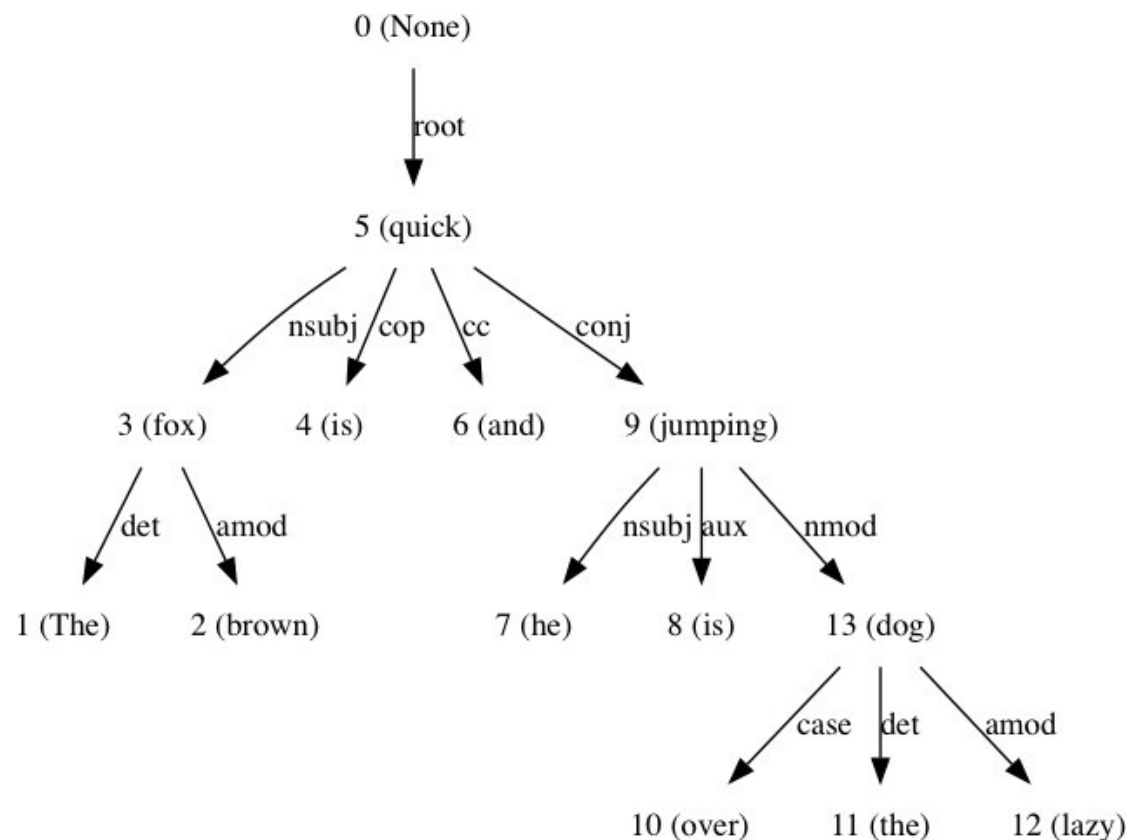
Each word has a **single head**

Dependency structure is **connected**

Projective

Acyclic

Unique path from ROOT to each word



3 - Summary




❖ 6 – Basic Vectorization

- ☐ One-hot encoding
- ☐ Bag-of-words (BoW)
- ☐ Bag-of-N-gram

Rome = [1, 0, 0, 0, 0, 0, ..., 0]
 Paris = [0, 1, 0, 0, 0, 0, ..., 0]
 Italy = [0, 0, 1, 0, 0, 0, ..., 0]
 France = [0, 0, 0, 1, 0, 0, ..., 0]



(Arrows indicate: Rome points to the first '1' in its vector, Paris points to the first '1' in its vector, and word V points to the last '0' in the first vector.)

TF*IDF weight vectors

	Set 1	Set 2
	0.023	0.000
	0.140	0.000
...
	0.000	0.010

3 - Summary


❖ 7 – Word2Vec



<i>man</i> →	0.6	-0.2	0.8	0.9	-0.1	-0.9	-0.7
<i>woman</i> →	0.7	0.3	0.9	-0.7	0.1	-0.5	-0.4
<i>king</i> →	0.5	-0.4	0.7	0.8	0.9	-0.7	-0.6
<i>queen</i> →	0.8	-0.1	0.8	-0.9	0.8	-0.5	-0.9

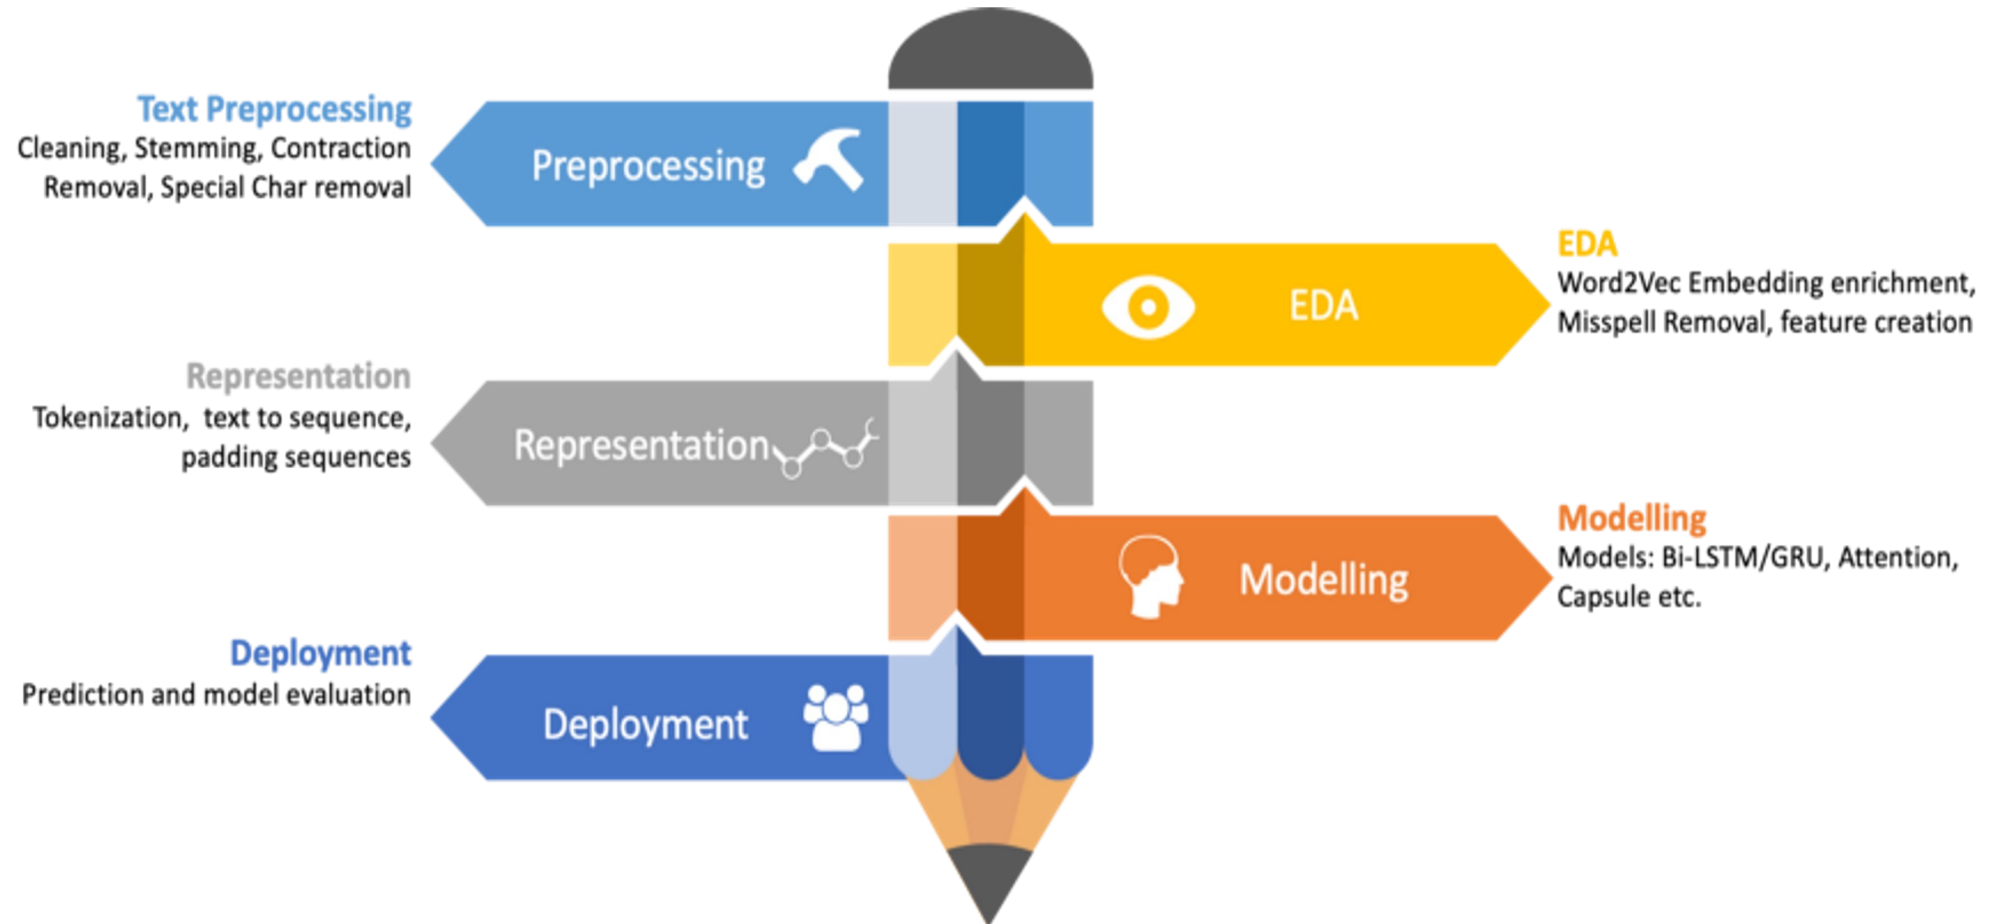
3 - Summary


❖ 7 – Pre-trained Embedding



3 - Summary

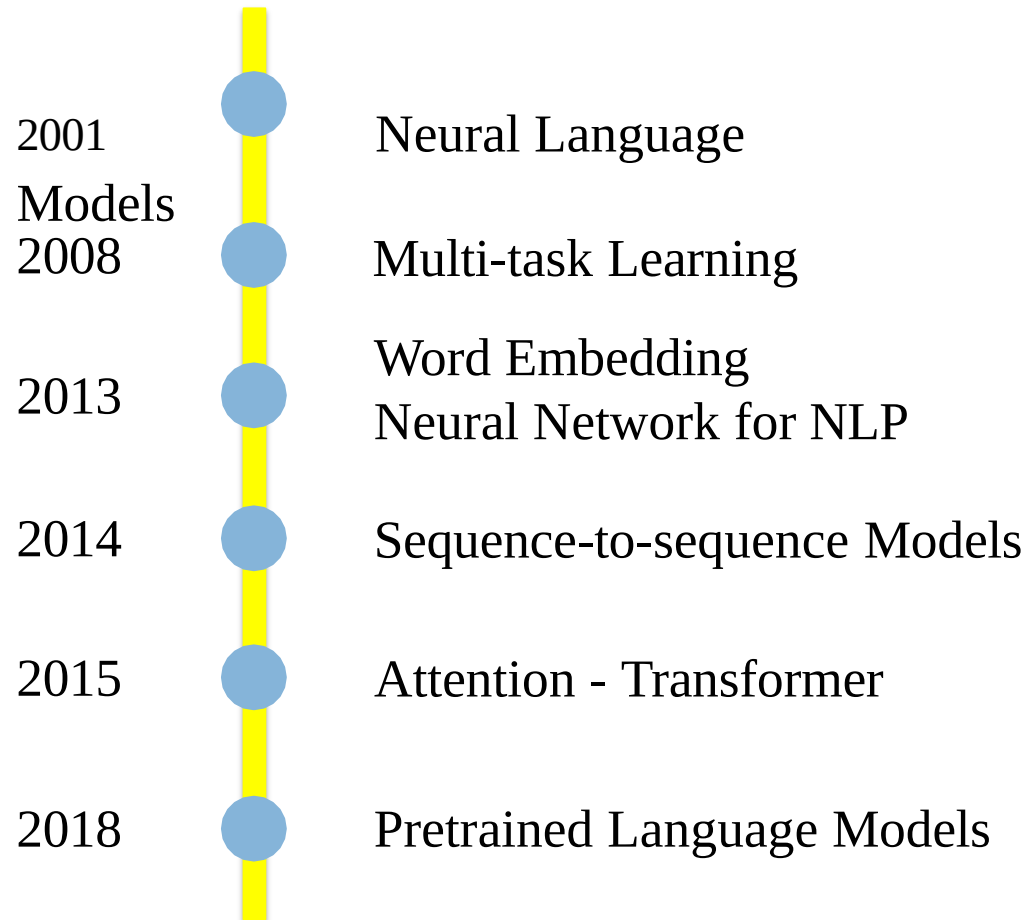
❖ NLP Pipeline





3 - Summary

❖ The neural history of NLP



3 - Summary

❖ The neural history of NLP

2001

Models

2008

2013

2014

2015

2018

Neural Language


Multi-task Learning

Word Embedding
Neural Network for NLP

Sequence-to-sequence Models

Attention - Transformer

Pretrained Language Models



3 - Summary

❖ The neural history of NLP

2001

Models

Neural Language

2008

2013

Word Embedding
Multi-task Learning
Neural Network for NLP

2014

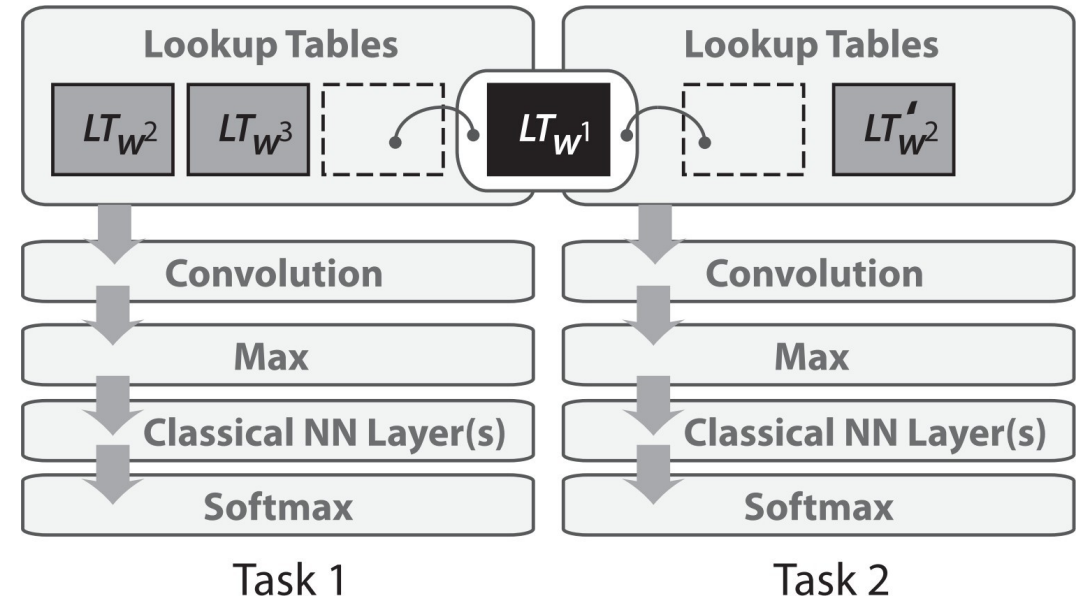
Sequence-to-sequence Models

2015

Attention - Transformer

2018

Pretrained Language Models



3 - Summary

❖ The neural history of NLP

2001
Models

Neural Language

2013

Word Embedding
Neural Network for NLP

2014

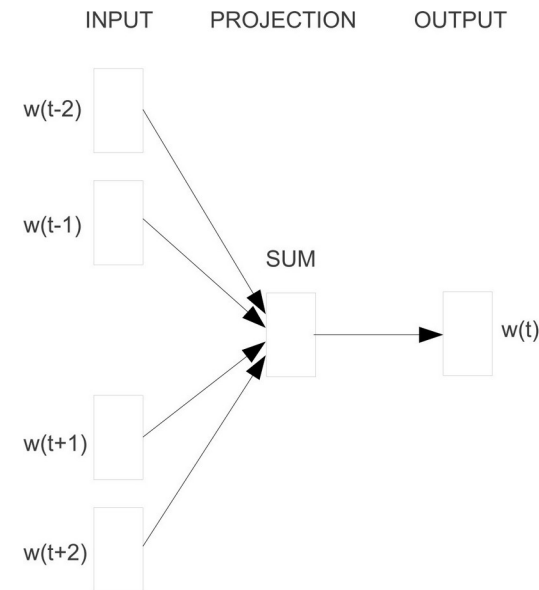
Sequence-to-sequence Models

2015


Attention - Transformer

2018

Pretrained Language Models



CBOW



Skip-gram

3 - Summary

❖ The neural history of NLP

2001
Models

Neural Language

2013

Word Embedding
Neural Network for NLP

2014

Sequence-to-sequence Models


2015

Attention - Transformer


2018

Pretrained Language Models


RNN




LSTM



GRU



CNN



3 - Summary

❖ The neural history of NLP

2001
Models
2008

Neural Language

Multi-task Learning

2013

Word Embedding
Neural Network for NLP

2014

Sequence-to-sequence Models

2015

Attention - Transformer

2018

Pretrained Language Models

Machine Translation


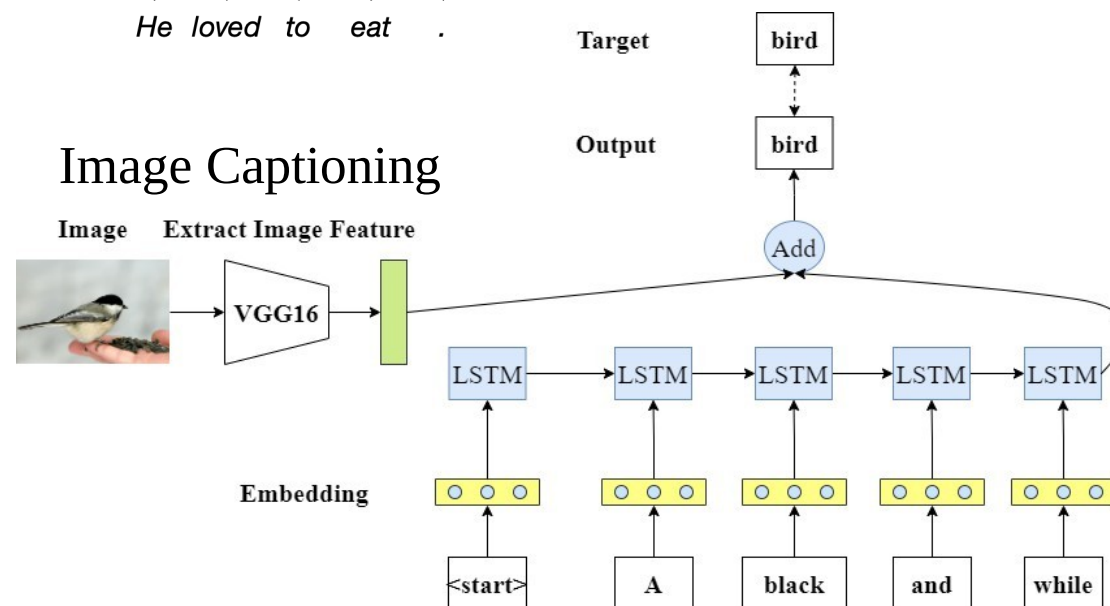


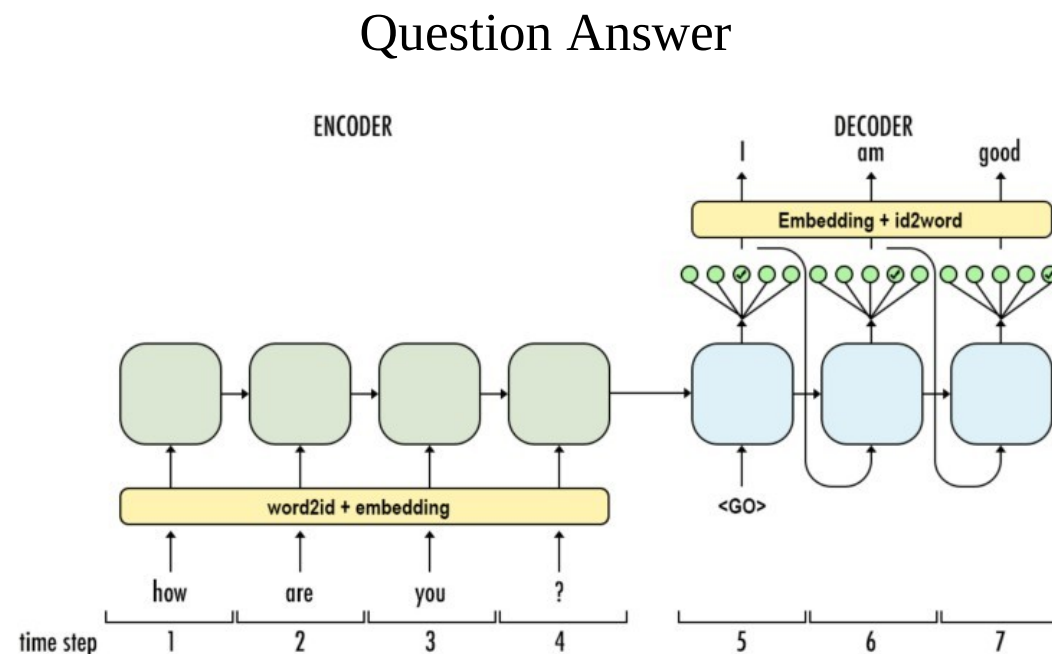
Image Captioning



3 - Summary

❖ The neural history of NLP

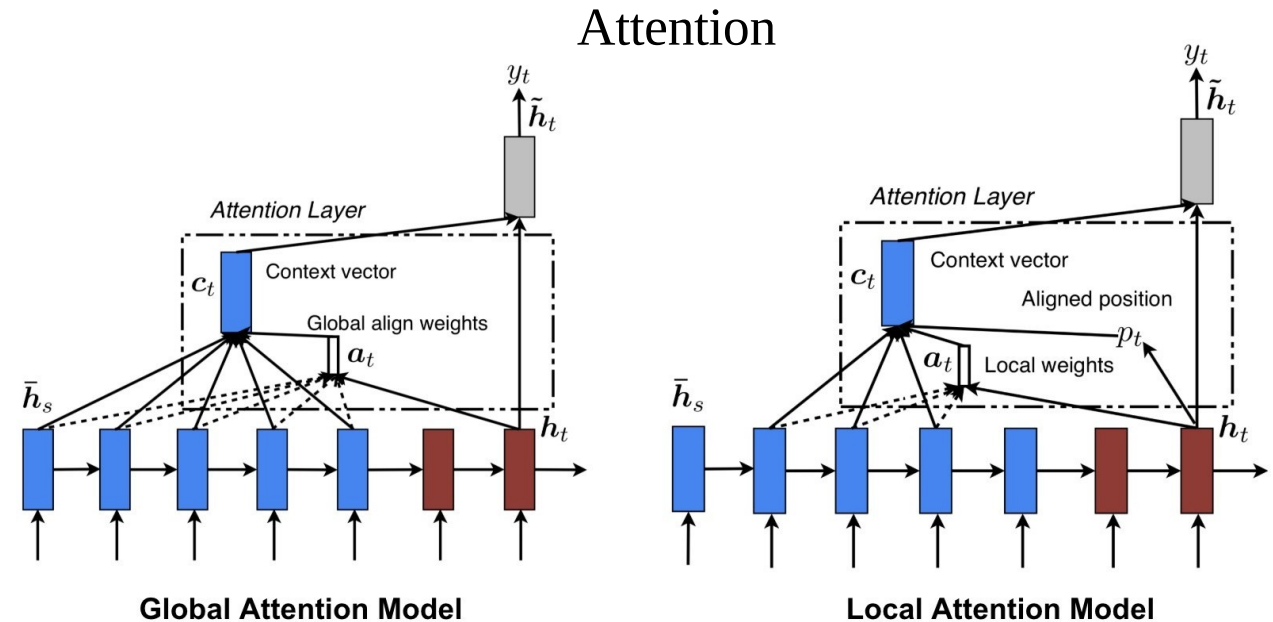
2001	Neural Language Models
2008	Multi-task Learning
2013	Word Embedding Neural Network for NLP
2014	Sequence-to-sequence Models
2015	Attention - Transformer
2018	Pretrained Language Models



3 - Summary

❖ The neural history of NLP

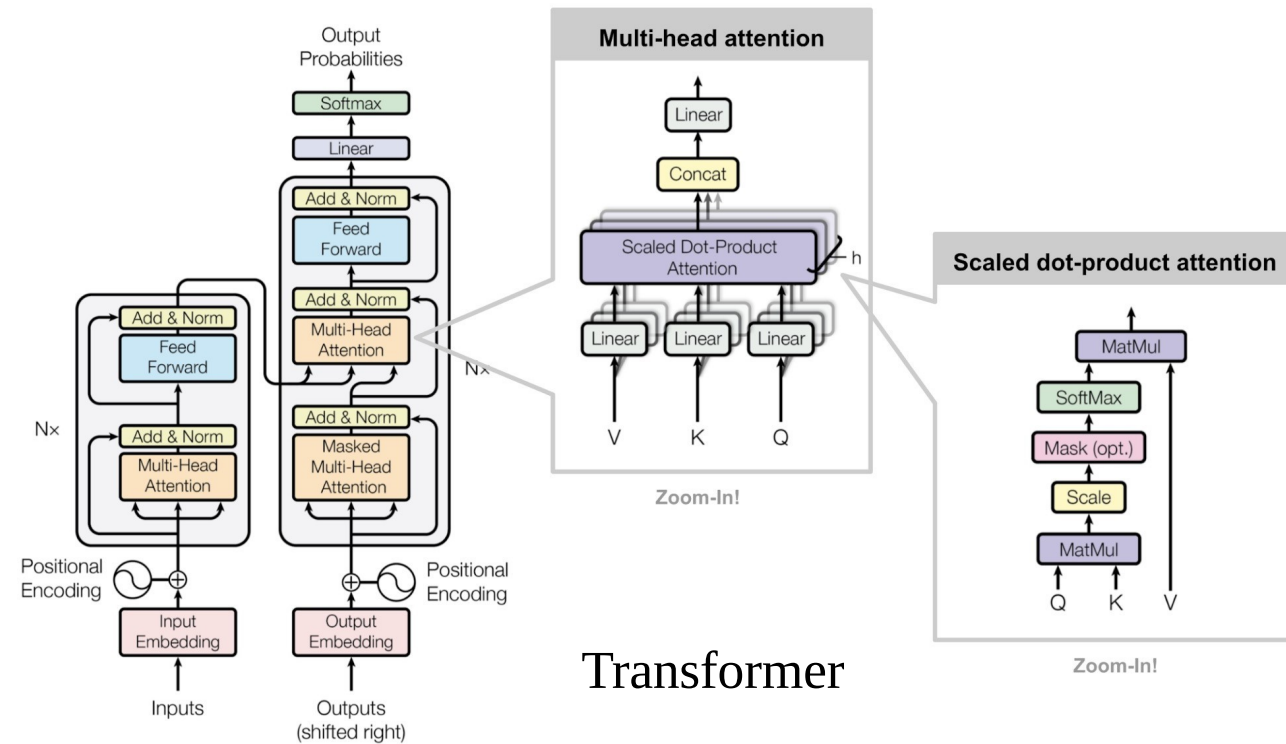
2001	Neural Language Models
2008	Multi-task Learning
2013	Word Embedding Neural Network for NLP
2014	Sequence-to-sequence Models
2015	Attention - Transformer
2018	Pretrained Language Models



3 - Summary

❖ The neural history of NLP

2001	Neural Language
Models	
2008	Multi-task Learning
2013	Word Embedding Neural Network for NLP
2014	Sequence-to-sequence Models
2015	Attention - Transformer
2018	Pretrained Language Models



Transformer

3 - Summary

❖ The neural history of NLP

2001

Models

2008

2013

2014

2015

2018

Neural Language


Multi-task Learning

Word Embedding
Neural Network for NLP

Sequence-to-sequence Models

Attention - Transformer

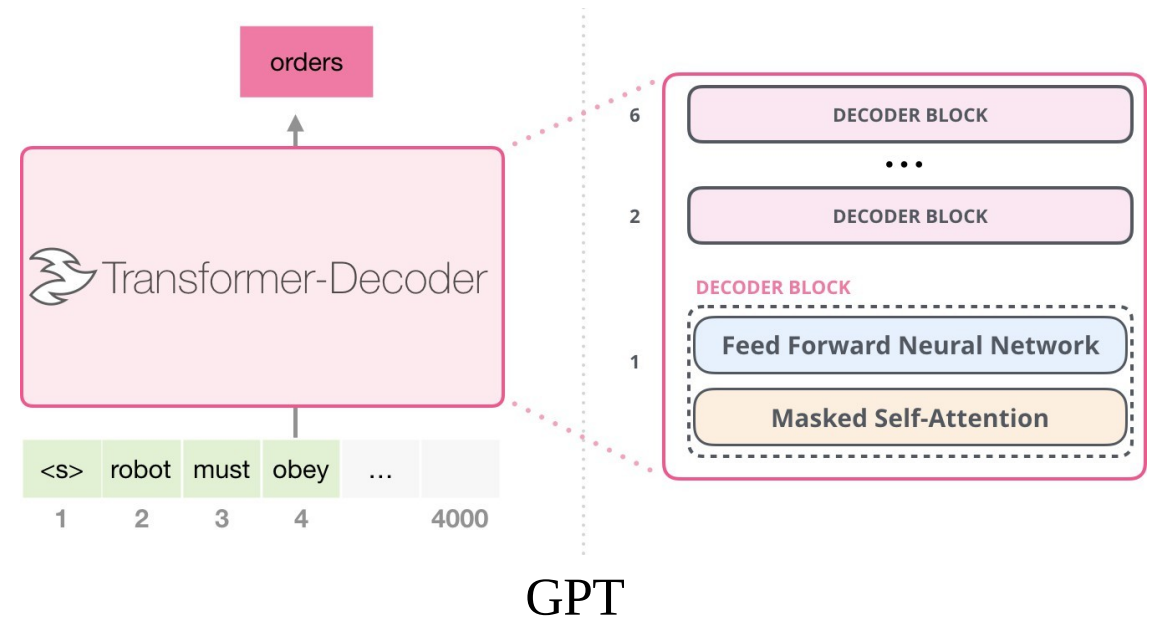
Pretrained Language Models



3 - Summary

❖ The neural history of NLP

2001	Neural Language
Models	
2008	Multi-task Learning
2013	Word Embedding Neural Network for NLP
2014	Sequence-to-sequence Models
2015	Attention - Transformer
2018	Pretrained Language Models



3 - Summary

❖ Reference

1. <https://web.stanford.edu/~jurafsky/slp3/>
2. <http://web.stanford.edu/class/cs224n/>
3. <https://d2l.ai/>
4. <http://nlpprogress.com/>
5. <https://github.com/undertheseanlp/NLP-Vietnamese-progress>



AI VIET NAM

@aivietnam.edu.vn

Thanks!

Any
questions?