



word2vec

From theory to practice

Hendrik Heuer

Stockholm NLP Meetup

About me

Hendrik Heuer
hendrikheuer@gmail.com
<http://hen-drik.de>
@hen_drik



“You shall know a word
by the company it keeps”

—J. R. Firth 1957

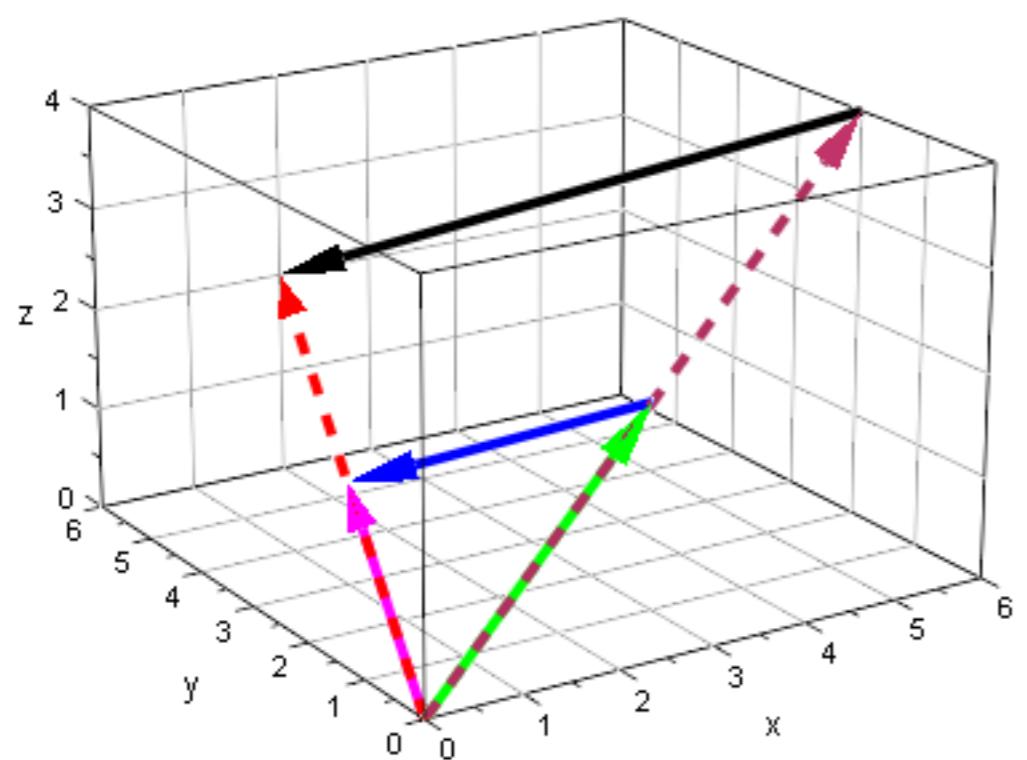
“You shall know a word
by the company it keeps”

One of the most successful ideas of modern statistical NLP

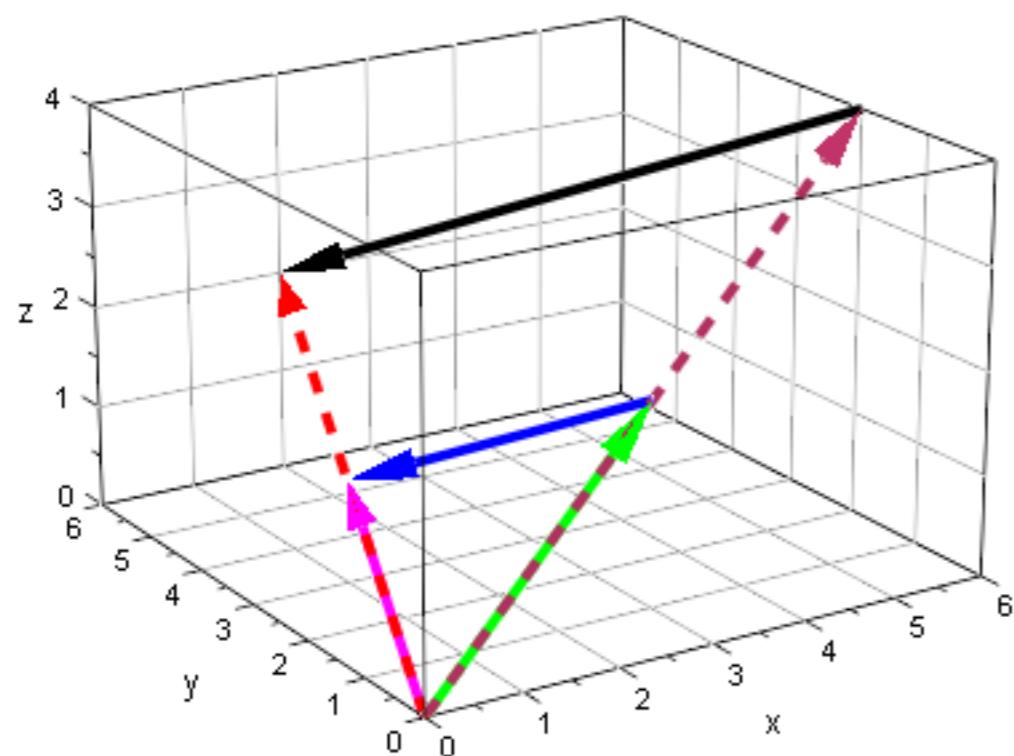
government debt problems turning into banking crises as has happened in

saying that Europe needs unified banking regulation to replace the hodgepodge

↖ These words will represent *banking* ↗



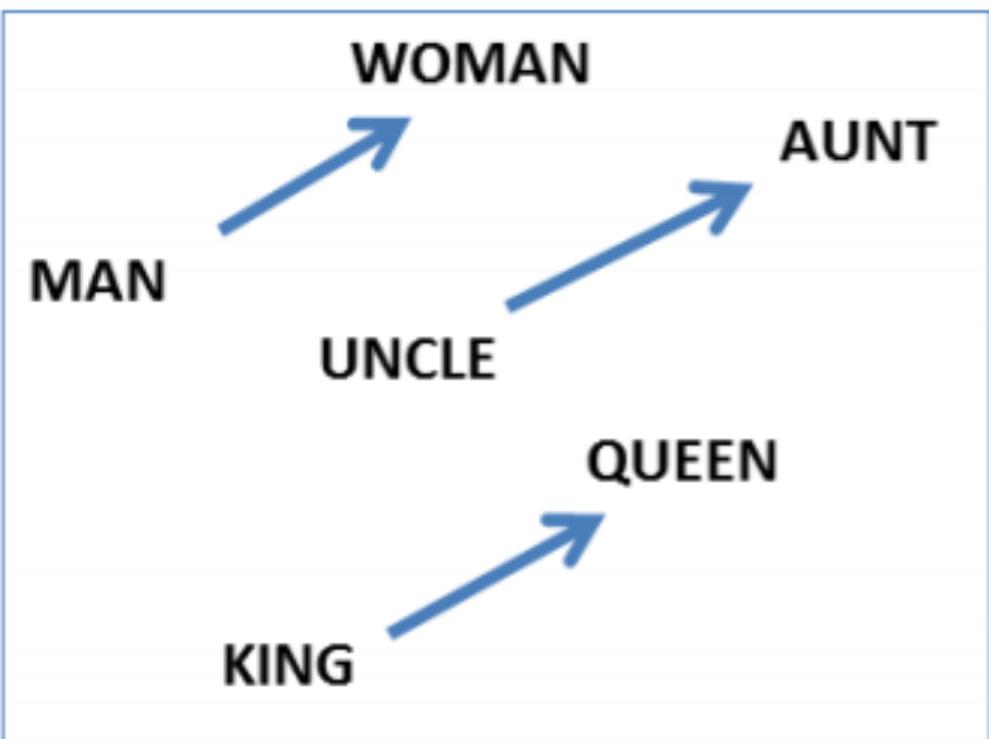
Quoted after Socher



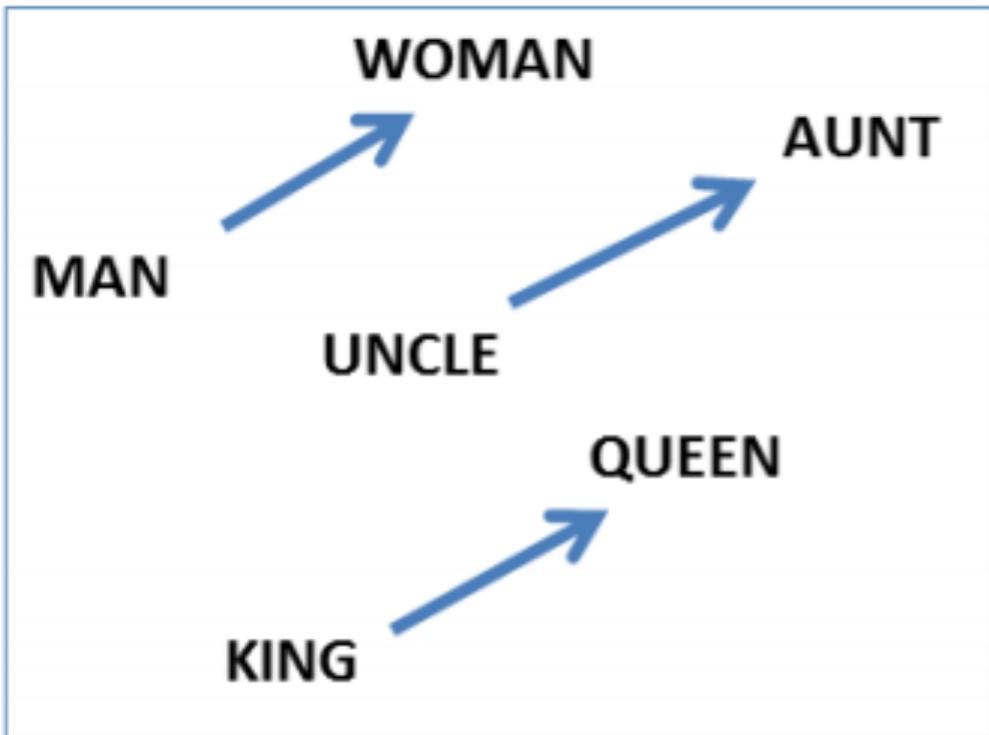
linguistics =

$$\begin{bmatrix} 0.286 \\ 0.792 \\ -0.177 \\ -0.107 \\ 0.109 \\ -0.542 \\ 0.349 \\ 0.271 \end{bmatrix}$$

Quoted after Socher

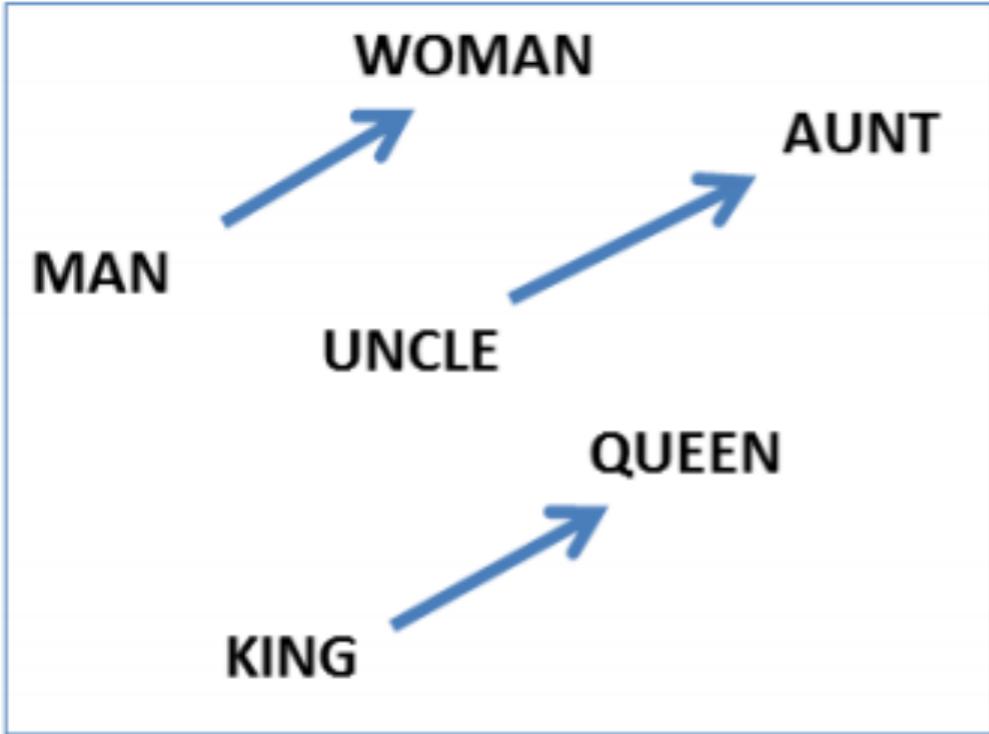


(Mikolov et al., NAACL HLT, 2013)



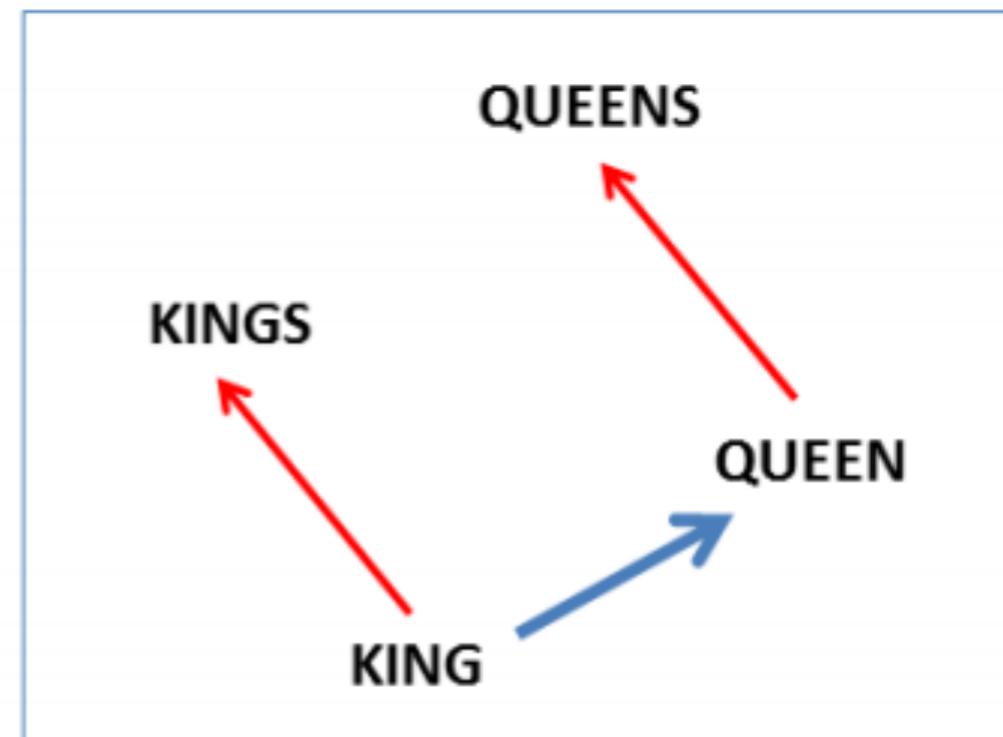
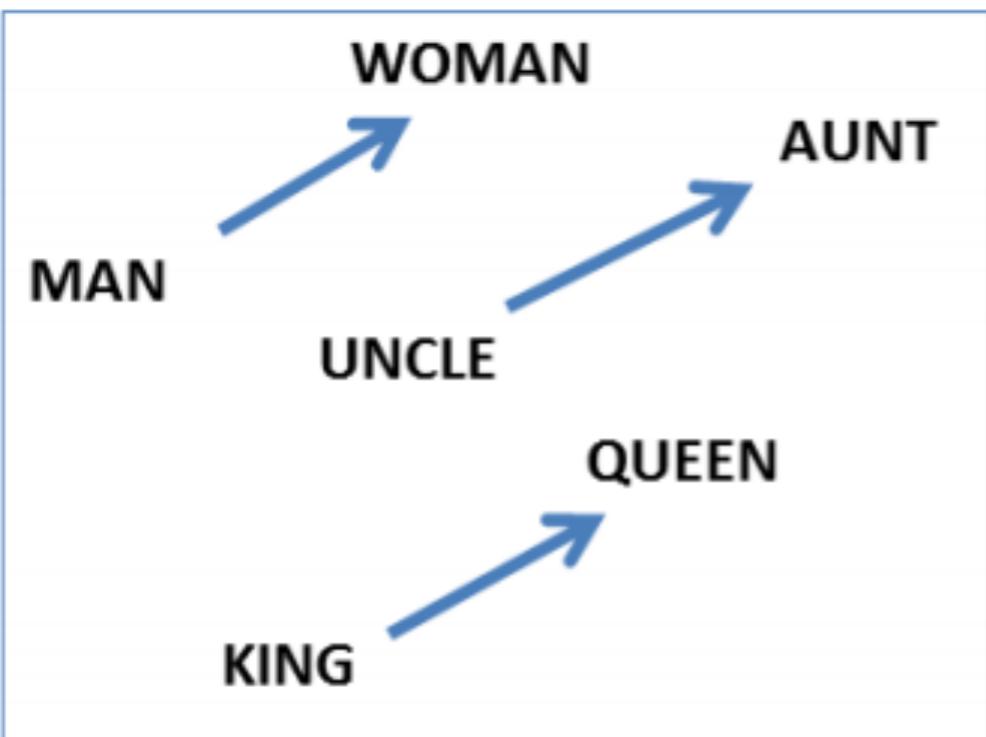
(Mikolov et al., NAACL HLT, 2013)

Vectors are directions in space
Vectors can encode relationships

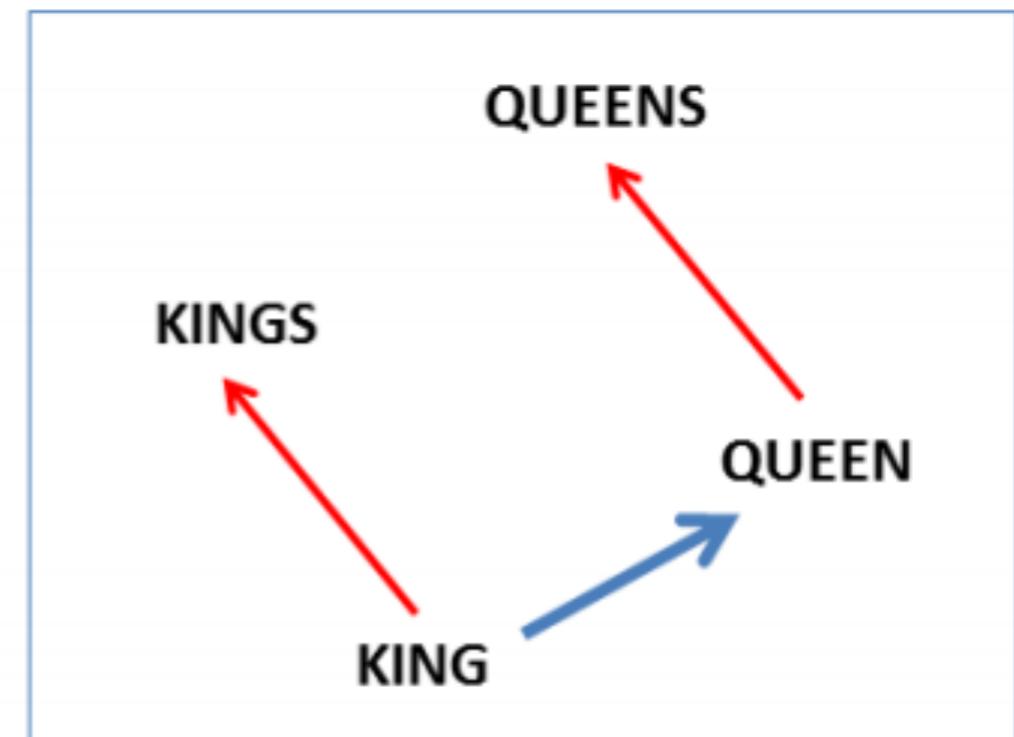
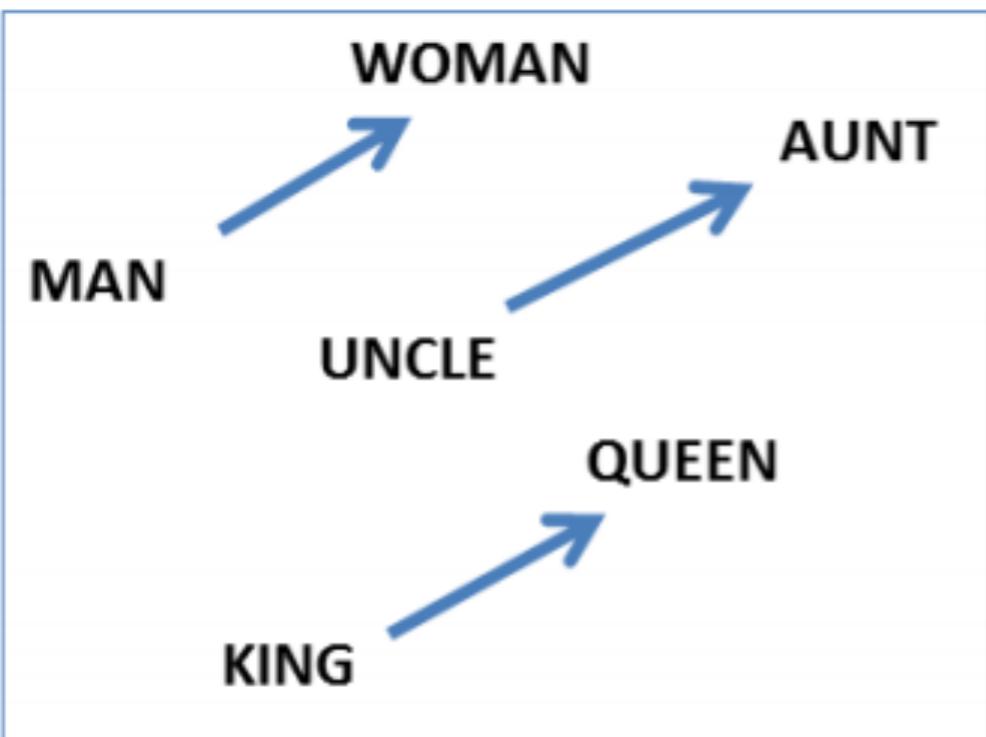


(Mikolov et al., NAACL HLT, 2013)

man is to **woman** as **king** is to ?



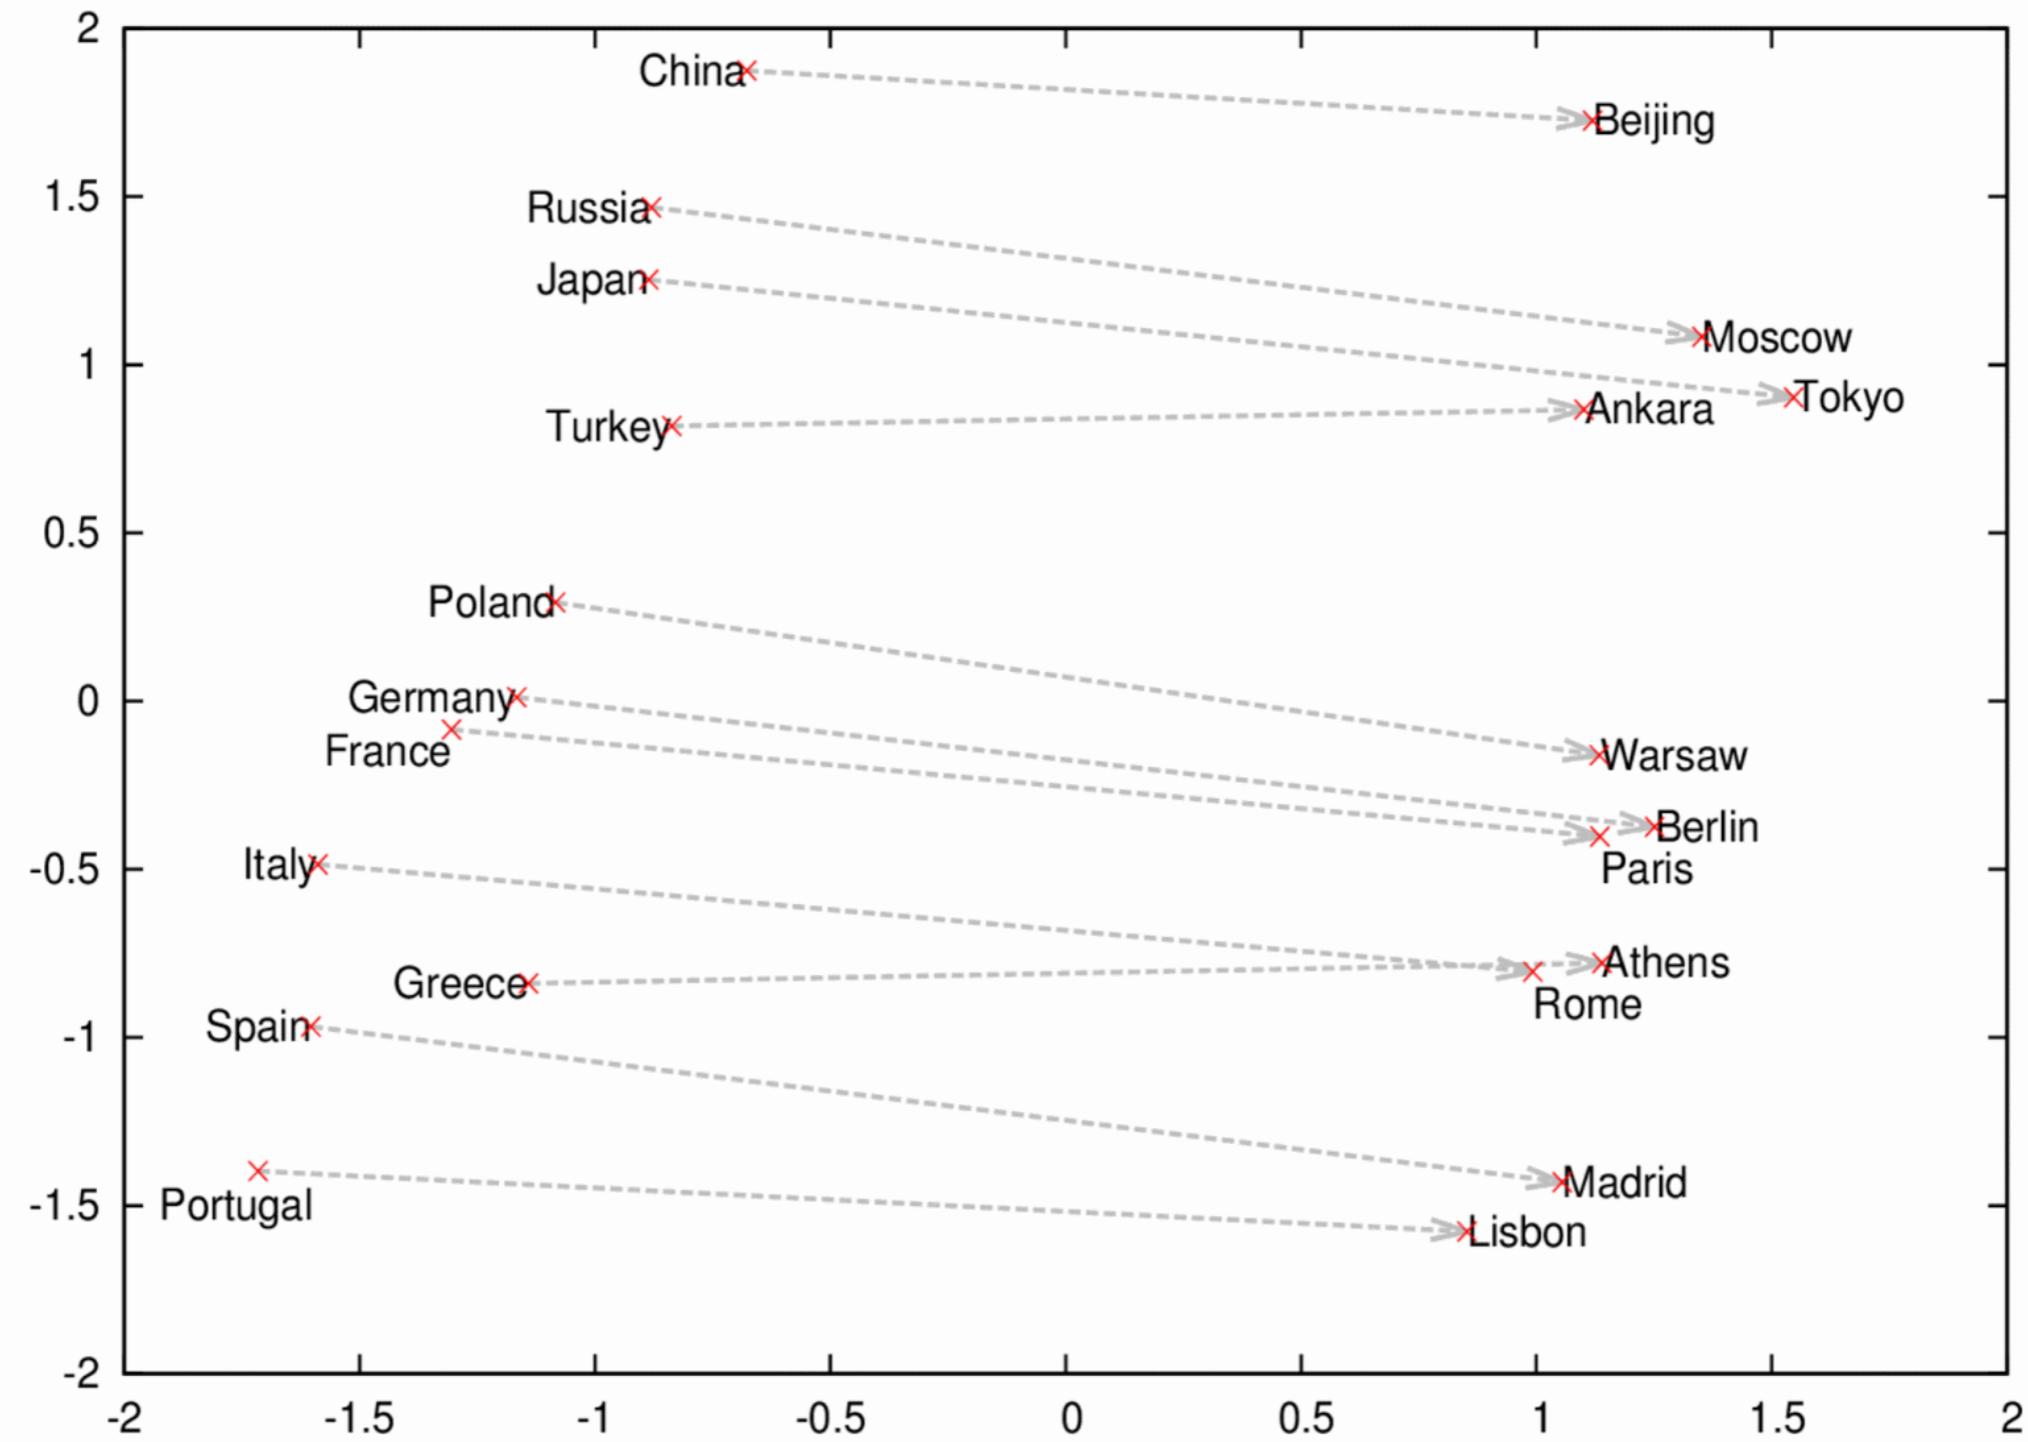
(Mikolov et al., NAACL HLT, 2013)



(Mikolov et al., NAACL HLT, 2013)



Country and Capital Vectors Projected by PCA



Efficient Estimation of Word Representations in Vector Space

Tomas Mikolov

Google Inc., Mountain View, CA

tmikolov@google.com

Kai Chen

Google Inc., Mountain View, CA

kaichen@google.com

Greg Corrado

Google Inc., Mountain View, CA

gcorrado@google.com

Jeffrey Dean

Google Inc., Mountain View, CA

jeff@google.com

Abstract

We propose two novel model architectures for computing continuous vector representations of words from very large data sets. The quality of these representations is measured in a word similarity task, and the results are compared to the previously best performing techniques based on different types of neural networks. We observe large improvements in accuracy at much lower computational cost, i.e. it takes less than a day to learn high quality word vectors from a 1.6 billion words data set. Furthermore, we show that these vectors provide state-of-the-art performance on our test set for measuring syntactic and semantic word similarities.

1 Introduction

Many current NLP systems and techniques treat words as atomic units - there is no notion of similarity between words, as these are represented as indices in a vocabulary. This choice has several good

word2vec

- by Mikolov, Sutskever, Chen, Corrado and Dean at Google
- NAACL 2013
- takes a **text corpus** as **input** and produces the **word vectors** as **output**

$$\text{linguistics} = \begin{pmatrix} 0.286 \\ 0.792 \\ -0.177 \\ -0.107 \\ 0.109 \\ -0.542 \\ 0.349 \\ 0.271 \end{pmatrix}$$

	Word	Cosine distance
Sweden		
Similar words		
	norway	0.760124
	denmark	0.715460
	finland	0.620022
	switzerland	0.588132
	belgium	0.585835
	netherlands	0.574631
	iceland	0.562368
	estonia	0.547621
	slovenia	0.531408
	floorball_federation	0.529570
	luxembourg	0.529477
	czech_republic	0.528778
	slovakia	0.526340
	romania	0.524281
	kista	0.522488
	helsinki_vantaa	0.519936
	swedish	0.519901
	balrog_ik	0.514556
	portugal	0.502495
	russia	0.500196
	slovakia_slovenia	0.496051
	ukraine	0.495712
	chur	0.484225
	thailand_togo	0.479172
	crimea_ukraine	0.478596
	panama_papua	0.477126
	latvia	0.477058

	Word	Cosine distance
Harvard		
Similar words		
	yale	0.638970
	cambridge	0.612665
	university	0.597709
	faculty	0.588422
	harvey_mudd	0.578338
	johns_hopkins	0.575645
	graduate	0.570294
	undergraduate	0.565881
	professor	0.563657
	mcgill	0.562168
	ph_d	0.558665
	california_berkeley	0.555539
	yale_university	0.550480
	harvard_crimson	0.549848
	princeton	0.544070
	college	0.542838
	oxford	0.531948
	barnard_college	0.530800
	professors	0.529959
	princeton_university	0.529763
	ucl	0.527395
	doctorates	0.526292
	doctoral	0.523317
	cambridge_massachusetts	0.519657
	juris_doctor	0.518845
	graduate_student	0.518815
	postgraduate	0.515757

word2vec

- **word meaning** and **relationships** between words are **encoded spatially**
- two main learning algorithms in word2vec:
continuous bag-of-words and
continuous skip-gram

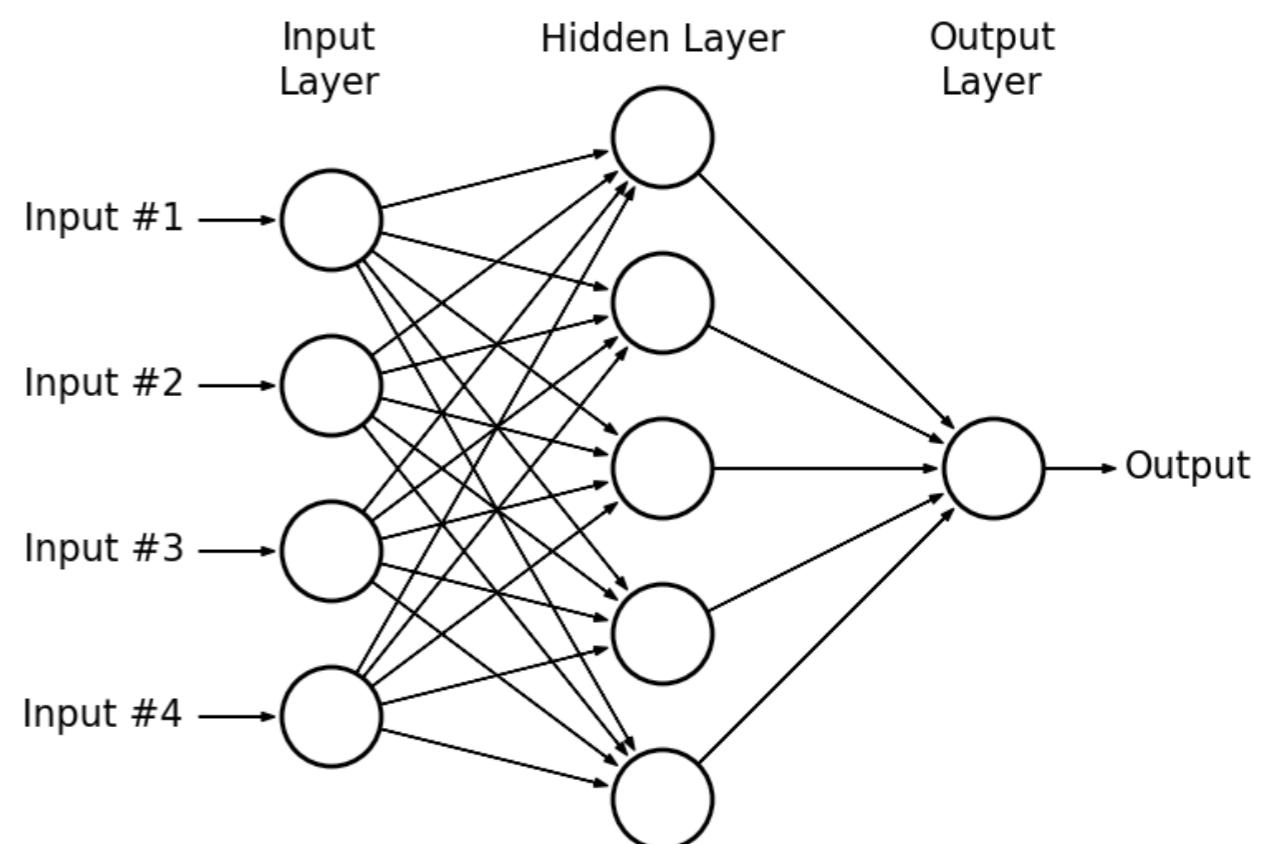
linguistics =

0.286
0.792
-0.177
-0.107
0.109
-0.542
0.349
0.271

Goal

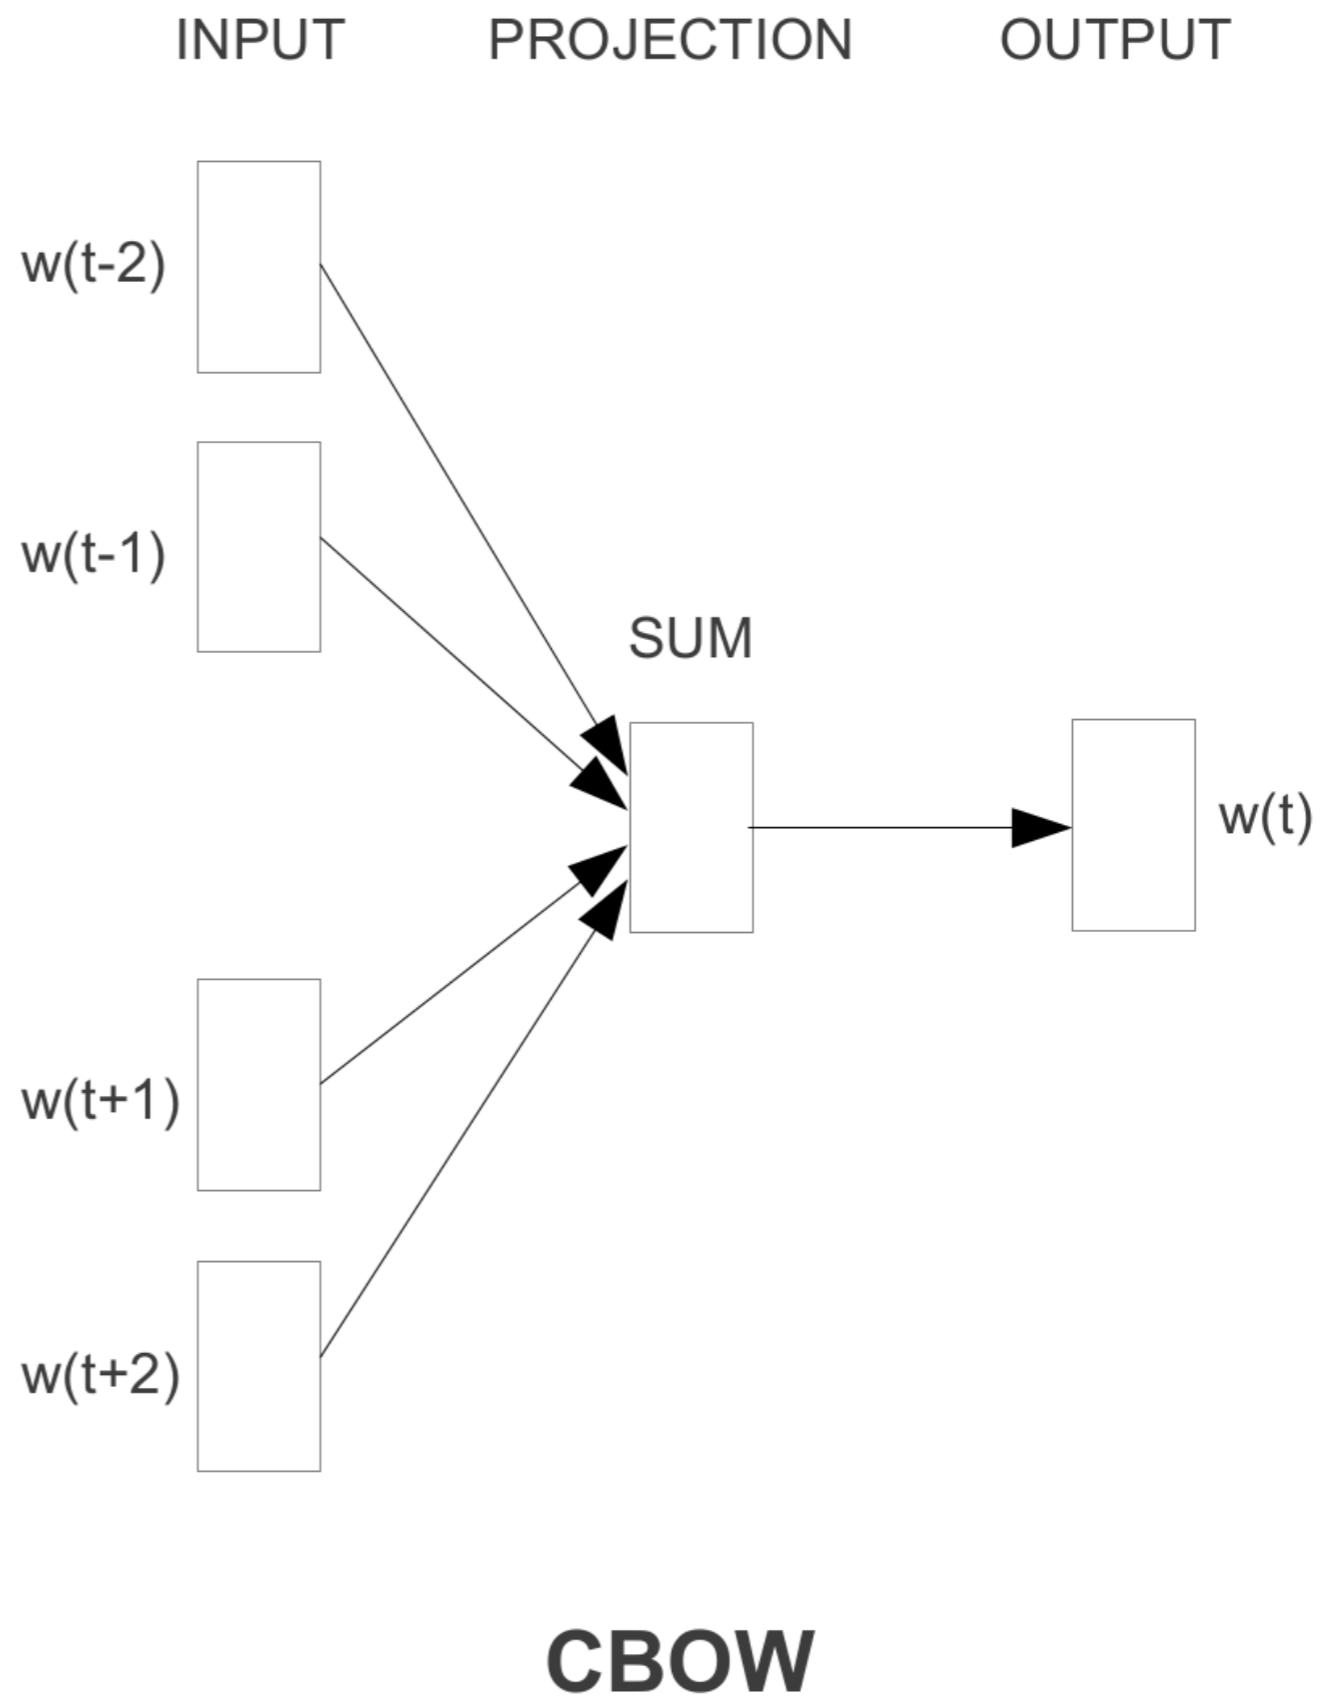
linguistics =

$$\begin{bmatrix} 0.286 \\ 0.792 \\ -0.177 \\ -0.107 \\ 0.109 \\ -0.542 \\ 0.349 \\ 0.271 \end{bmatrix}$$



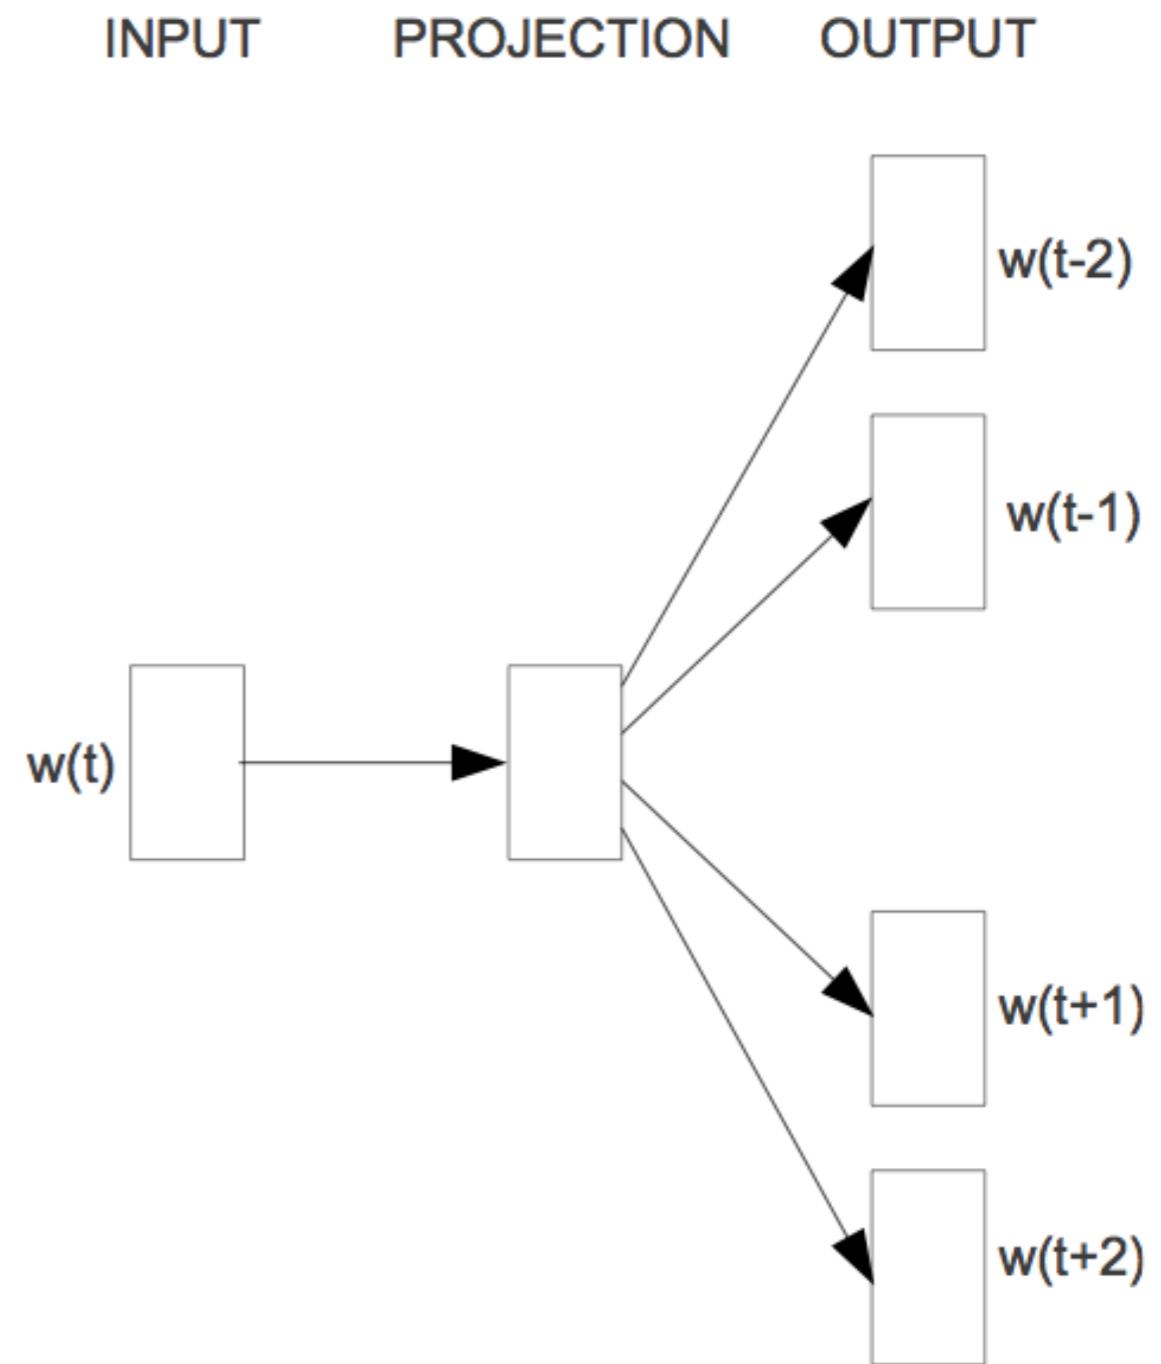
continuous bag-of-words

- predicting the current word **based on the context**
- **order of words in the history does not influence the projection**
- faster & **more appropriate for larger corpora**



continuous skip-gram

- maximize classification of a word **based on another word in the same sentence**
- better word vectors for **frequent words**, but slower to train



Skip-gram

Why it is awesome

- there is a **fast open-source implementation**
- can be used as **features** for **natural language processing tasks** and **machine learning** algorithms

Machine Translation

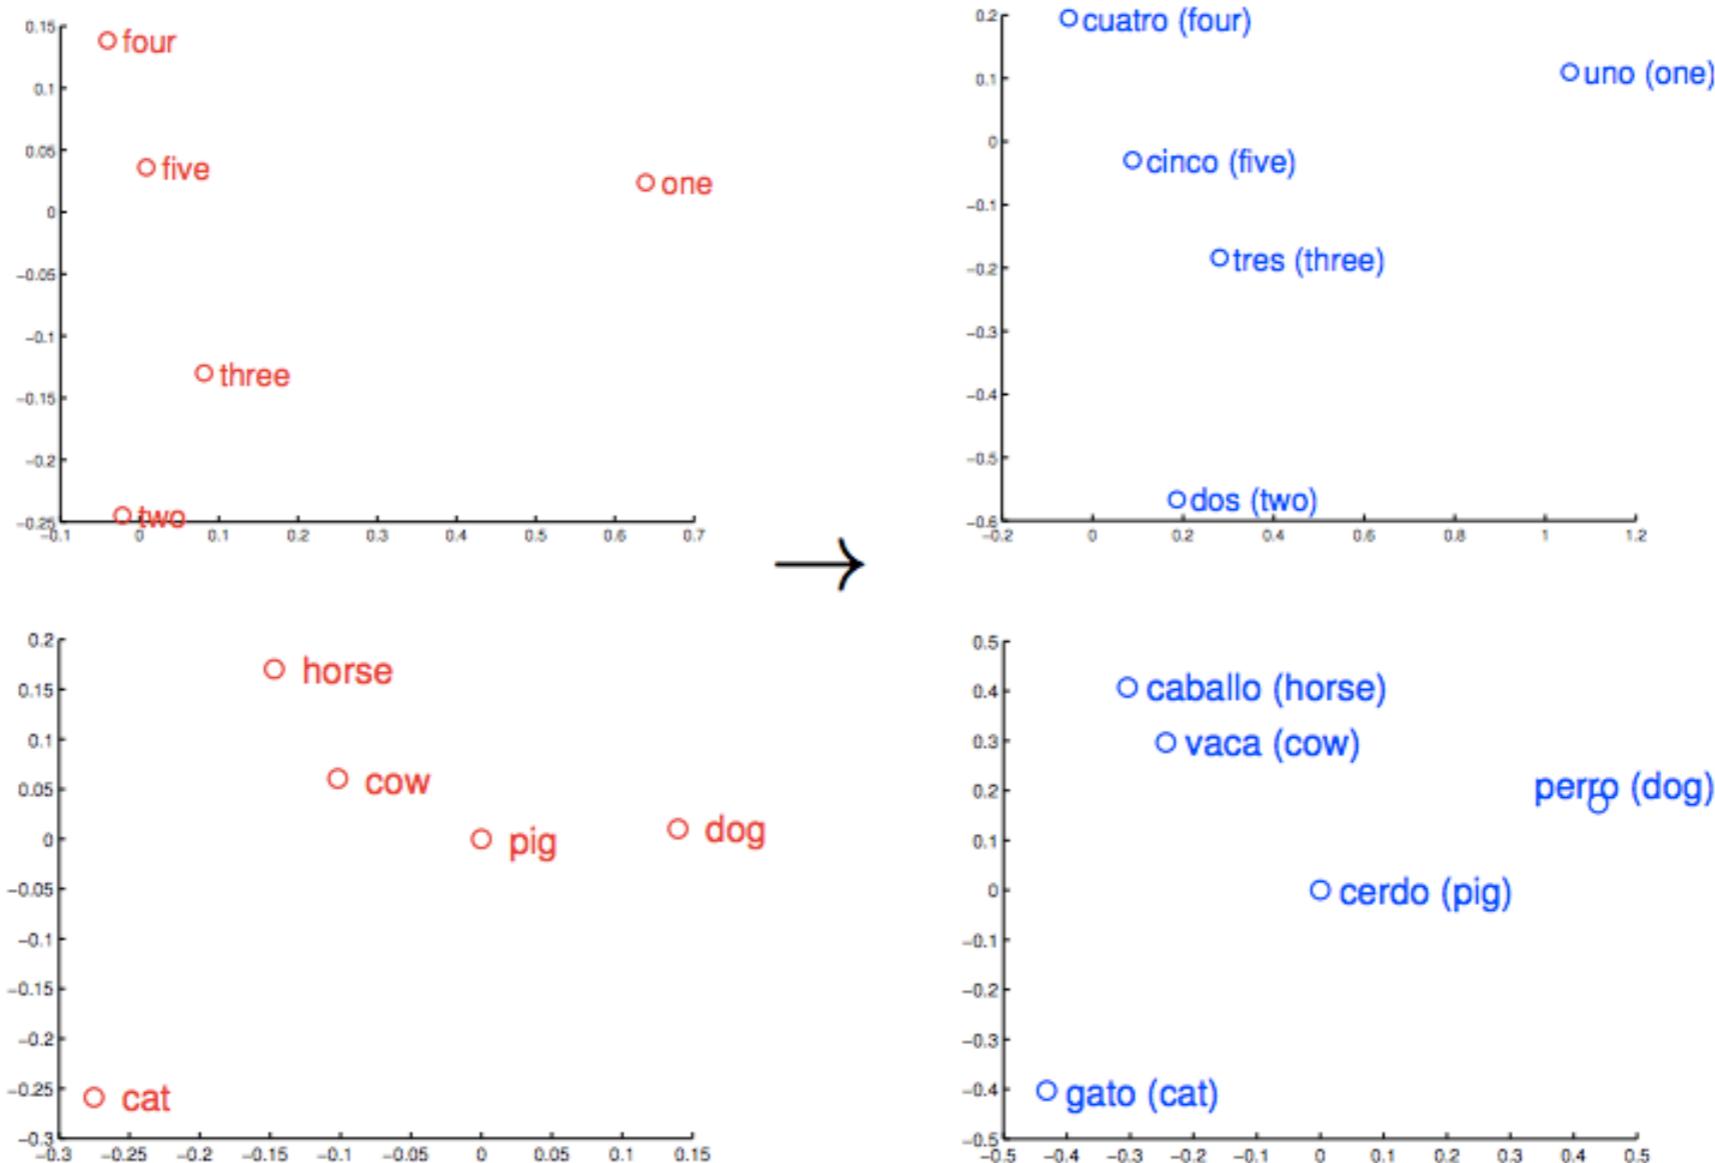


Figure 1: Distributed word vector representations of numbers and animals in English (left) and Spanish (right). The five vectors in each language were projected down to two dimensions using PCA, and then manually rotated to accentuate their similarity. It can be seen that these concepts have similar geometric arrangements in both spaces, suggesting that it is possible to learn an accurate linear mapping from one space to another. This is the key idea behind our method of translation.

Quoted after Mikolov

Sentiment Analysis

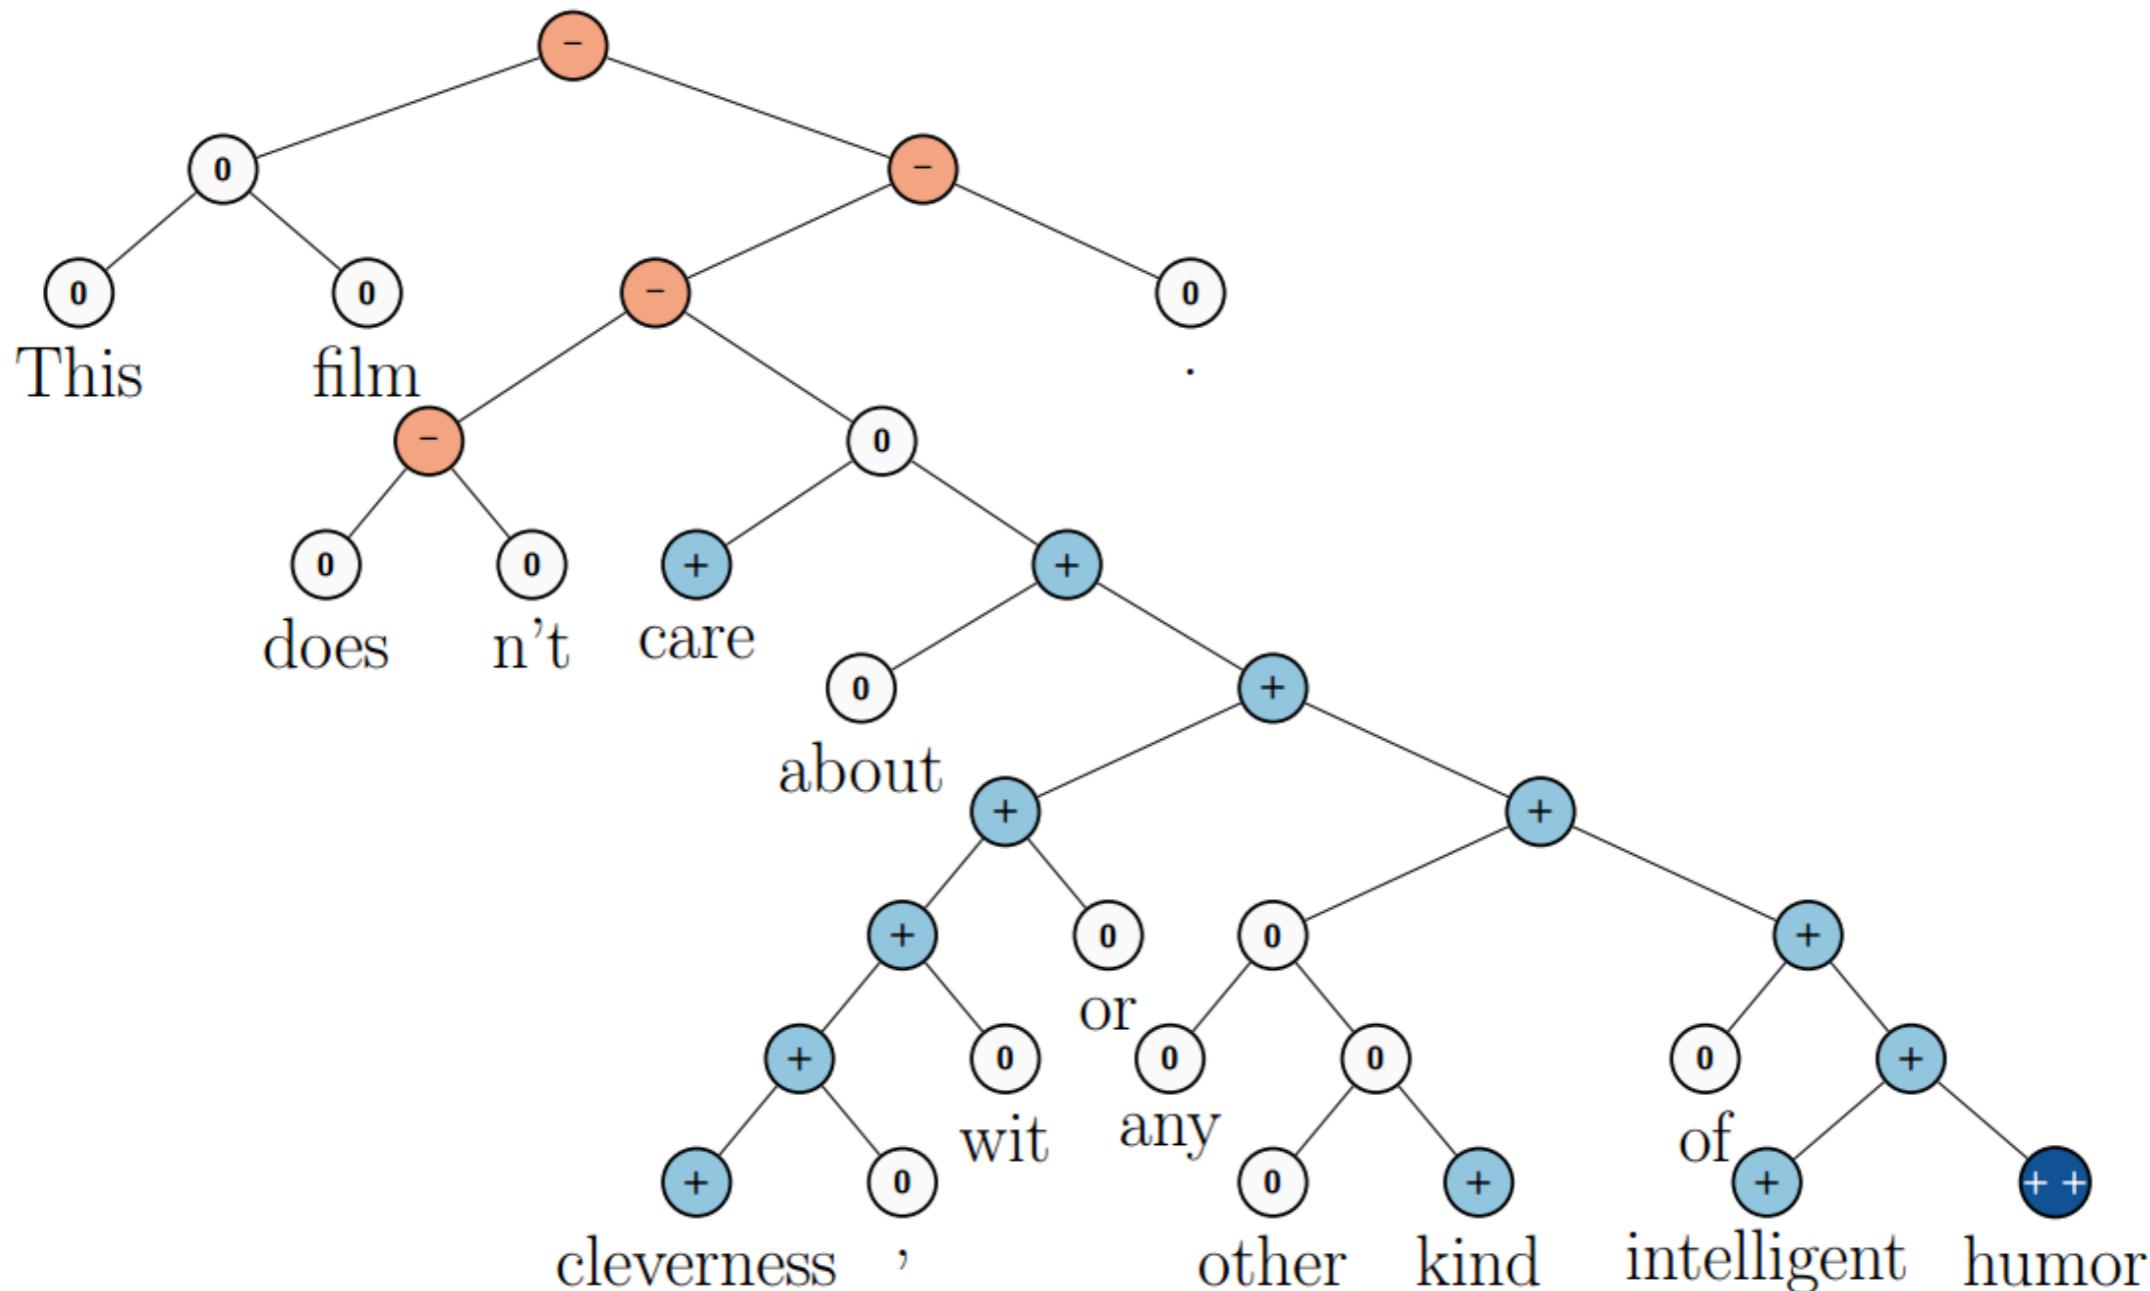


Image Descriptions

A person riding a motorcycle on a dirt road.



A group of young people playing a game of frisbee.



A herd of elephants walking across a dry grass field.



Two dogs play in the grass.



Two hockey players are fighting over the puck.



A close up of a cat laying on a couch.



A skateboarder does a trick on a ramp.



A little girl in a pink hat is blowing bubbles.



A red motorcycle parked on the side of the road.



A dog is jumping to catch a frisbee.



A refrigerator filled with lots of food and drinks.



A yellow school bus parked in a parking lot.



Describes without errors

Describes with minor errors

Somewhat related to the image

Unrelated to the image

Figure 5. A selection of evaluation results, grouped by human rating.

Quoted after Vinyals

Using word2vec

- Original: <http://word2vec.googlecode.com/svn/trunk/>
- C++11 version: <https://github.com/jdeng/word2vec>
- Python: <http://radimrehurek.com/gensim/models/word2vec.html>
- Java: https://github.com/ansjsun/word2vec_java
- Parallel java: <https://github.com/siegfang/word2vec>
- CUDAversion: <https://github.com/whatupbiatch/cuda-word2vec>

Using it in Python





gensim

topic modelling for humans



Download

latest version from the Python Package Index



Direct install with:
easy_install -U gensim

[Home](#)[Tutorials](#)[Install](#)[Support](#)[API](#)[About](#)

```
>>> from gensim import corpora, models, similarities  
>>>  
>>> # Load corpus iterator from a Matrix Market file on disk.  
>>> corpus = corpora.MmCorpus('/path/to/corpus.mm')  
>>>  
>>> # Initialize Latent Semantic Indexing with 200 dimensions.  
>>> lsi = models.LsiModel(corpus, num_topics=200)  
>>>  
>>> # Convert another corpus to the latent space and index it.  
>>> index = similarities.MatrixSimilarity(lsi[another_corpus])  
>>>  
>>> # Compute similarity of a query vs. indexed documents  
>>> sims = index[query]
```

Gensim is a FREE Python library



Scalable statistical semantics



Analyze plain-text documents for semantic structure



Retrieve semantically similar documents

Features

Hover your mouse over each feature for more info.



Scalability



Platform independent



Robust



Open source



Efficient implementations



Converters & I/O formats



Similarity queries



Support

Usage

```
model = word2vec.Word2Vec.load_word2vec_format(word2vec_vectors_filename, binary=True)

for word in words:
    if word in model:
        print model[word]
```

Training a model

```
1 >>> # import modules & set up logging
2 >>> import gensim, logging
3 >>> logging.basicConfig(format='%(asctime)s : %(levelname)s : %(message)s', level=logging.INFO)
4 >>>
5 >>> sentences = [['first', 'sentence'], ['second', 'sentence']]
6 >>> # train word2vec on the two sentences
7 >>> model = gensim.models.Word2Vec(sentences, min_count=1)
```

Training a model with iterator

```
1  >>> class MySentences(object):
2  ...     def __init__(self, dirname):
3  ...         self.dirname = dirname
4  ...
5  ...     def __iter__(self):
6  ...         for fname in os.listdir(self.dirname):
7  ...             for line in open(os.path.join(self.dirname, fname)):
8  ...                 yield line.split()
9  ...
10    >>> sentences = MySentences('/some/directory') # a memory-friendly iterator
11    >>> model = gensim.models.Word2Vec(sentences)
```

Doing it in C

- Download the code:
`git clone https://github.com/h10r/word2vec-macosx-maverics.git`
- Run 'make' to compile word2vec tool
- Run the demo scripts:

`./demo-word.sh`

and

`./demo-phrases.sh`

Doing it in C

- Download the code:
git clone **https://github.com/h10r/word2vec-macosx-maverics.git**

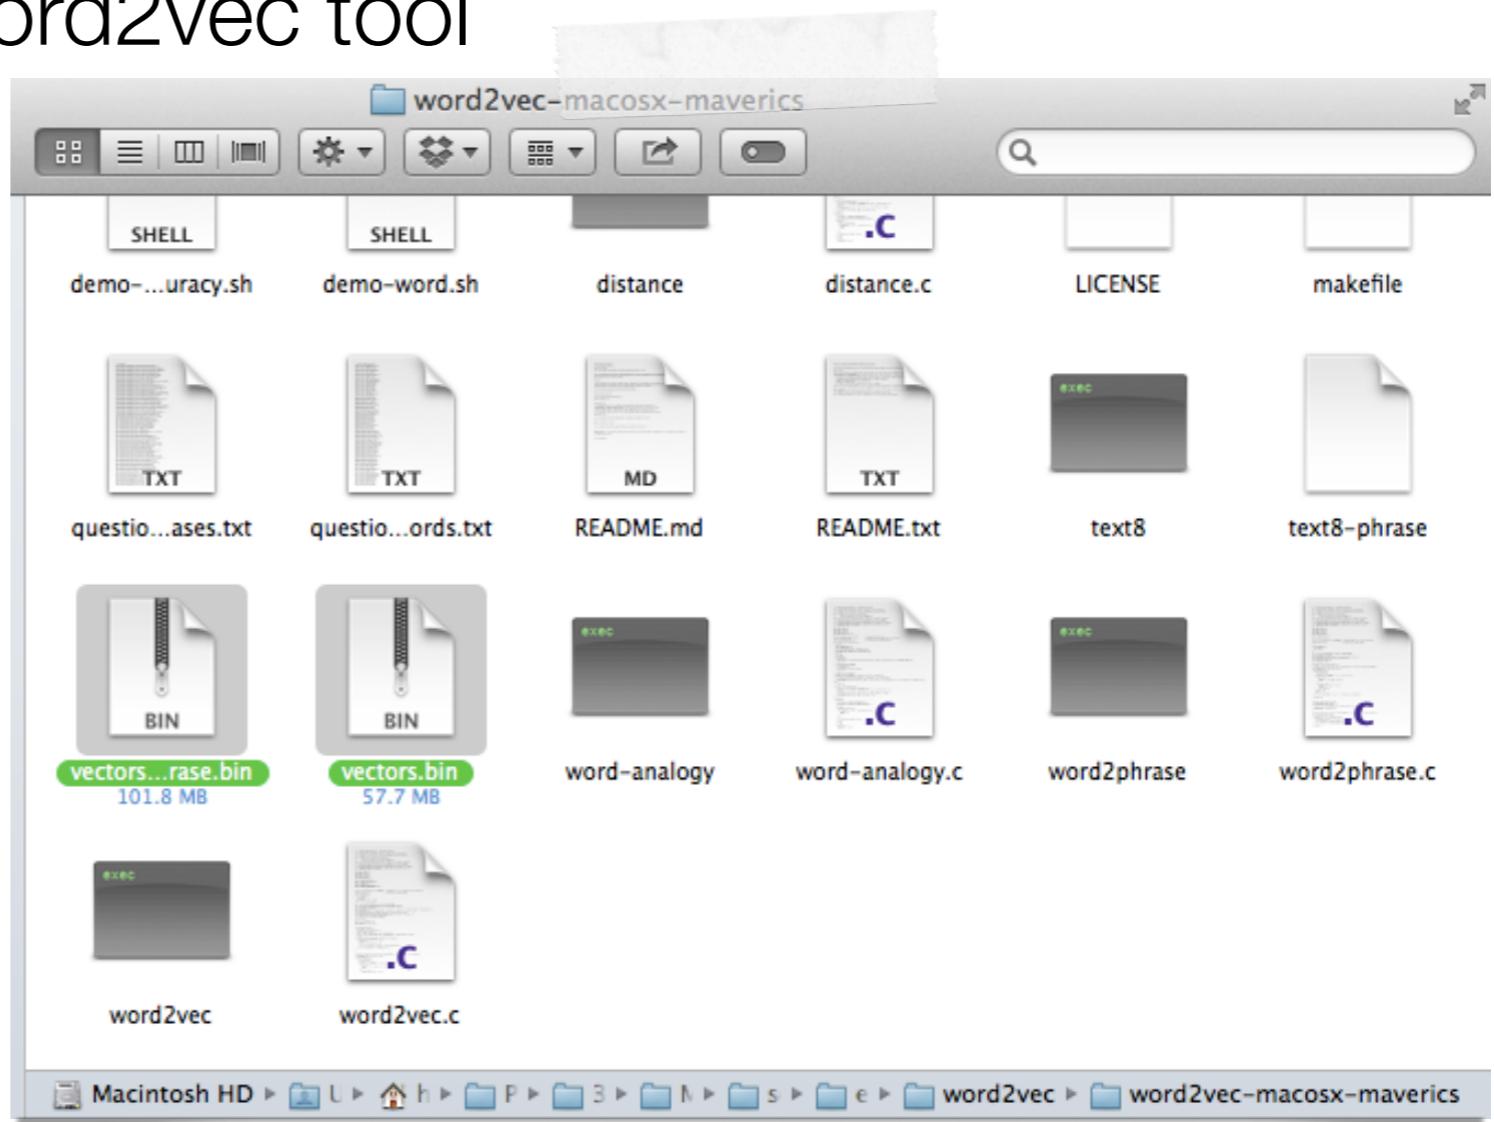
- Run 'make' to compile word2vec tool

- Run the demo scripts:

./demo-word.sh

and

./demo-phrases.sh



Pre-trained entity vectors with Freebase naming

We are also offering more than 1.4M pre-trained entity vectors with naming from [Freebase](#). This is especially helpful for projects related to knowledge mining.

- Entity vectors trained on 100B words from various news articles: [freebase-vectors-skipgram1000.bin.gz](#)
- Entity vectors trained on 100B words from various news articles, using the deprecated /en/ naming (more easily readable); the vectors are sorted by frequency: [freebase-vectors-skipgram1000-en.bin.gz](#)

Here is an example output of `./distance freebase-vectors-skipgram1000-en.bin`:

```
Enter word or sentence (EXIT to break): /en/geoffrey_hinton
```

Word	Cosine distance

/en/marvin_minsky	0.457204
/en/paul_corkum	0.443342
/en/william_richard_peltier	0.432396
/en/brenda_milner	0.430886
/en/john_charles_polanyi	0.419538
/en/leslie_valiant	0.416399

Testing a model

Google have released their testing set of about 20,000 syntactic and semantic test examples, following the "A is to B as C is to D" task:

<http://word2vec.googlecode.com/svn/trunk/questions-words.txt>.

Gensim support the same evaluation set, in exactly the same format:

```
1  >>> model.accuracy('/tmp/questions-words.txt')
2  2014-02-01 22:14:28,387 : INFO : family: 88.9% (304/342)
3  2014-02-01 22:29:24,006 : INFO : gram1-adjective-to-adverb: 32.4% (263/812)
4  2014-02-01 22:36:26,528 : INFO : gram2-opposite: 50.3% (191/380)
5  2014-02-01 23:00:52,406 : INFO : gram3-comparative: 91.7% (1222/1332)
6  2014-02-01 23:13:48,243 : INFO : gram4-superlative: 87.9% (617/702)
7  2014-02-01 23:29:52,268 : INFO : gram5-present-participle: 79.4% (691/870)
8  2014-02-01 23:57:04,965 : INFO : gram7-past-tense: 67.1% (995/1482)
9  2014-02-02 00:15:18,525 : INFO : gram8-plural: 89.6% (889/992)
10 2014-02-02 00:28:18,140 : INFO : gram9-plural-verbs: 68.7% (482/702)
11 2014-02-02 00:28:18,140 : INFO : total: 74.3% (5654/7614)
```

nutrition

neurological

addiction

disorders obesity

diet

respond

disorder

pathological craving

alcohol

purchase

sales
employees marketing

stores

profits

corporate

reporting

advertising

expenses

advertisements

bulletin consumers loyalty

customer

email

laboratory

experimental

experiments

experiment

studies

researchers scientists

integrated
components

technology
computing

embedded
systems

architectures

audio

tools

distributed
unix

implementations

platform

ntation

users

software

default

microsoft

user

linux

wikipedia

ver

apis

cal

toolkit

scripts

download

apache

eval

https

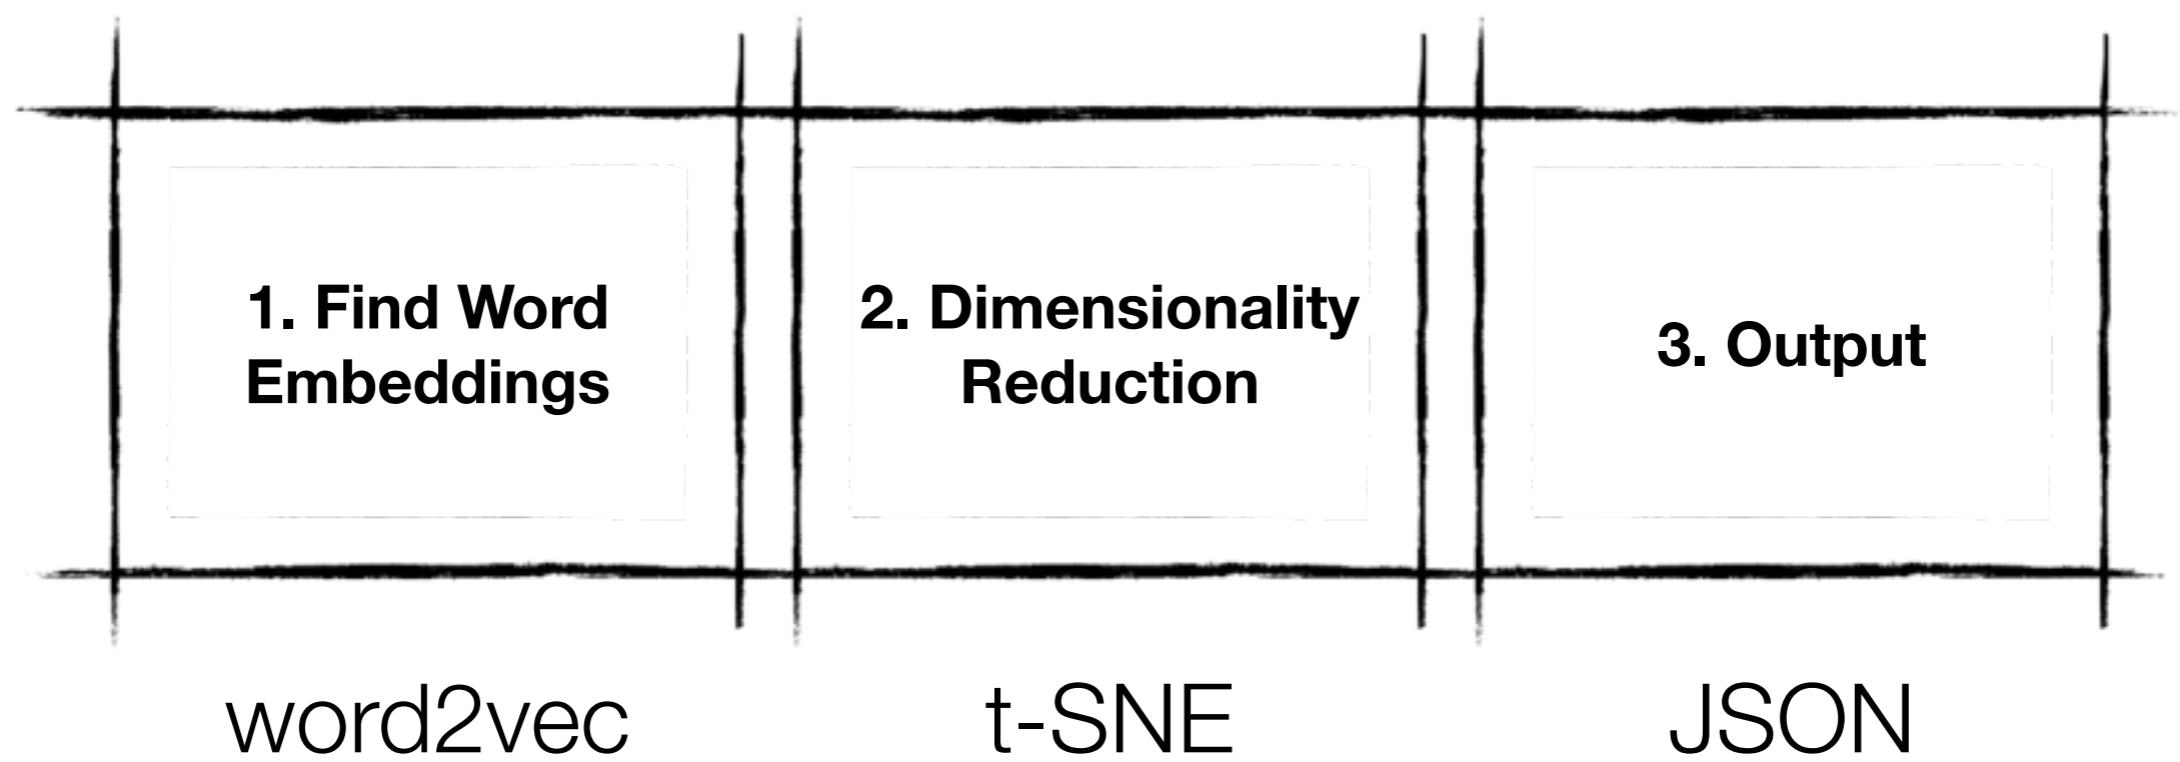
google

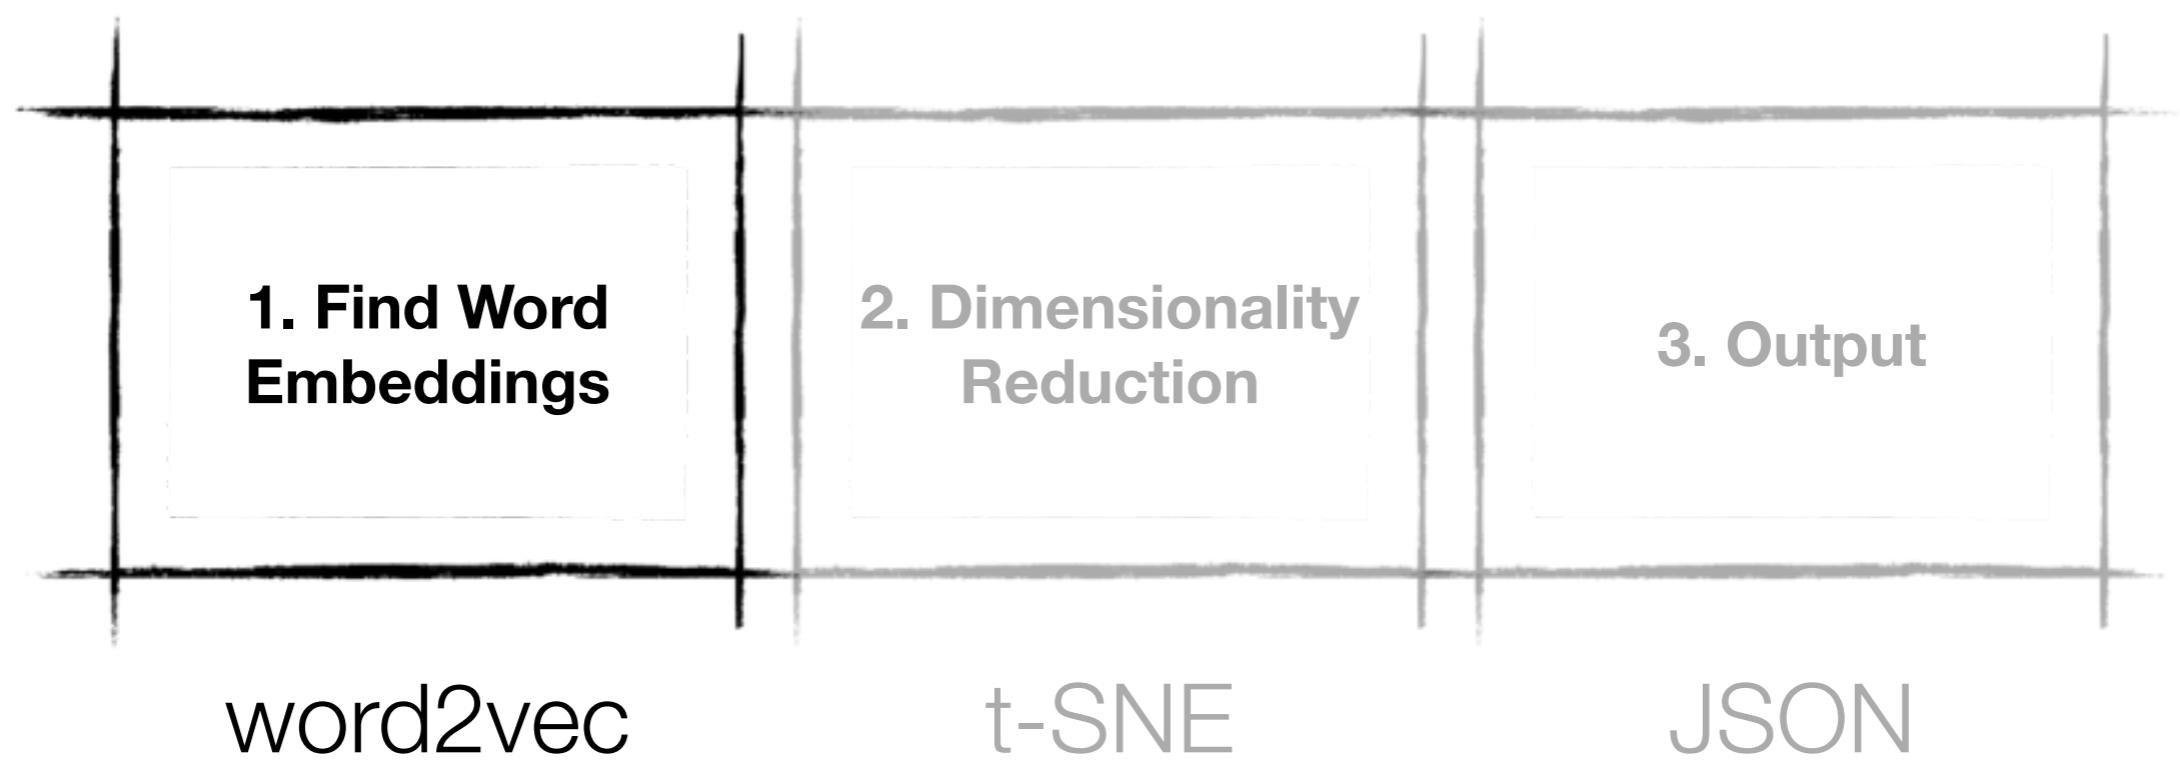
docum

word2vec & t-SNE

+ Zoom in - Zoom out Toggle - Topics A Toggle - Topics B Toggle - Topics Intersection







```
def parse_file_with_gensim( filename ):
    model = word2vec.Word2Vec.load_word2vec_format( "data/vectors.bin", binary=True )

    frequent_words = load_list_of_frequent_words()

    keep_track_of_words = set()

    word_vectors = []
    word_labels = []

    f = codecs.open( filename, encoding='utf-8' )
    for line in f:
        line = line.lower()
        line = " ".join(["" if ch in string.punctuation else ch for ch in line]).split()

        words = nltk.word_tokenize( line )

        for word in words:
            if word not in keep_track_of_words:
                if not word in frequent_words:
                    if word in model:
                        word_vectors.append( model[word] )
                        word_labels.append( word )

                keep_track_of_words.add( word )

    return word_vectors, word_labels
```

```
def parse_file_with_gensim( filename ):
    model = word2vec.Word2Vec.load_word2vec_format( "data/vectors.bin", binary=True )

    frequent_words = load_list_of_frequent_words()

    keep_track_of_words = set()

    word_vectors = []
    word_labels = []

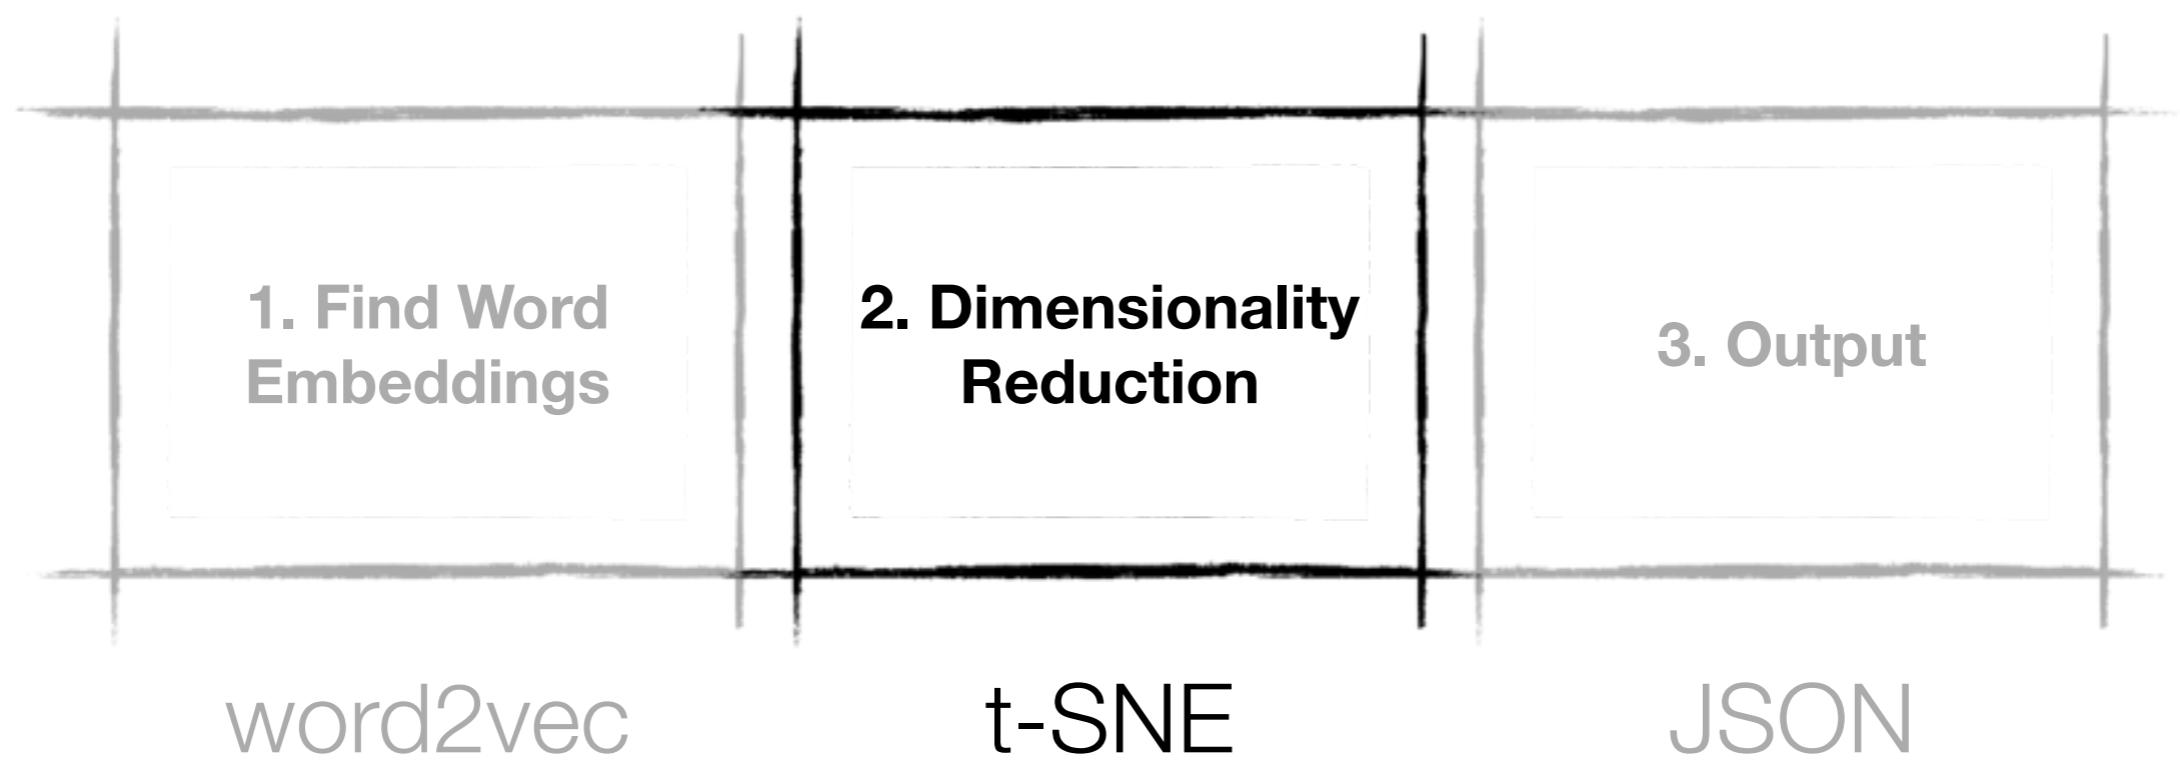
    f = codecs.open( filename, encoding='utf-8' )
    for line in f:
        line = line.lower()
        line = " ".join([" " if ch in string.punctuation else ch for ch in line]).split()

        words = nltk.word_tokenize( line )

        for word in words:
            if word not in keep_track_of_words:
                if not word in frequent_words:
                    if word in model:
                        word_vectors.append( model[word] )
                        word_labels.append( word )

            keep_track_of_words.add( word )

    return word_vectors, word_labels
```



Visualizing Data using t-SNE

Laurens van der Maaten

L.VANDERMAATEN@MICC.UNIMAAS.NL

MICC-IKAT

Maastricht University

P.O. Box 616, 6200 MD Maastricht, The Netherlands

Geoffrey Hinton

HINTON@CS.TORONTO.EDU

Department of Computer Science

University of Toronto

6 King's College Road, M5S 3G4 Toronto, ON, Canada

Editor: Leslie Pack Kaelbling

Abstract

We present a new technique called “t-SNE” that visualizes high-dimensional data by giving each datapoint a location in a two or three-dimensional map. The technique is a variation of Stochastic Neighbor Embedding (Hinton and Roweis, 2002) that is much easier to optimize, and produces significantly better visualizations by reducing the tendency to crowd points together in the center of the map. t-SNE is better than existing techniques at creating a single map that reveals structure at many different scales. This is particularly important for high-dimensional data that lie on several different, but related, low-dimensional manifolds, such as images of objects from multiple classes seen from multiple viewpoints. For visualizing the structure of very large datasets, we show how t-SNE can use random walks on neighborhood graphs to allow the implicit structure of all of the data to influence the way in which a subset of the data is displayed. We illustrate the performance



Previous: [sklearn.manifold...
sklearn.manifold...](#) Next: [sklearn.manifold...
sklearn.manifold...](#) Up: [Reference](#)

This documentation is for
scikit-learn version
0.15.2 — Other versions

If you use the software,
please consider citing
[scikit-learn](#).

[sklearn.manifold.TSNE](#)

sklearn.manifold.TSNE

```
class sklearn.manifold.TSNE(n_components=2, perplexity=30.0, early_exaggeration=4.0, learning_rate=1000.0, n_iter=1000, metric='euclidean', init='random', verbose=0, random_state=None)
```

t-distributed Stochastic Neighbor Embedding.

t-SNE [1] is a tool to visualize high-dimensional data. It converts similarities between data points to joint probabilities and tries to minimize the Kullback-Leibler divergence between the joint probabilities of the low-dimensional embedding and the high-dimensional data. t-SNE has a cost function that is not convex, i.e. with different initializations we can get different results.

It is highly recommended to use another dimensionality reduction method (e.g. PCA for dense data or TruncatedSVD for sparse data) to reduce the number of dimensions to a reasonable amount (e.g. 50) if the number of features is very high. This will suppress some noise and speed up the computation of pairwise distances between samples. For more tips see Laurens van der Maaten's FAQ [2].

Parameters: `n_components` : int, optional (default: 2)

Dimension of the embedded space.

`perplexity` : float, optional (default: 30)

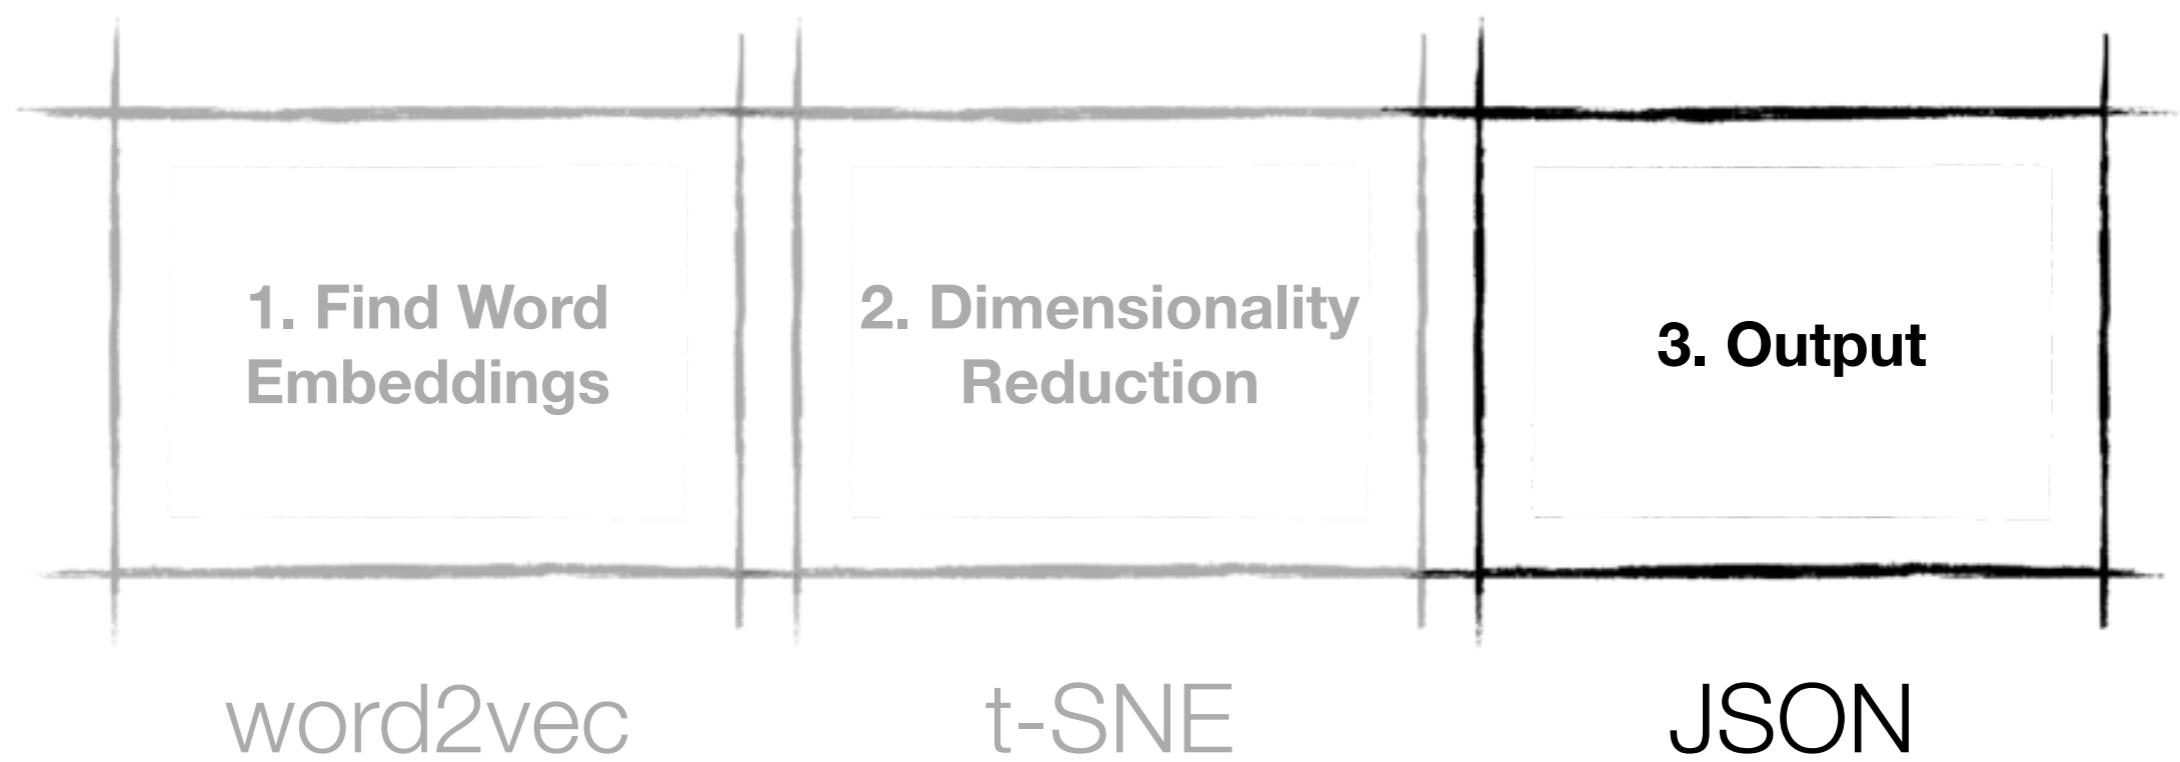
The perplexity is related to the number of nearest neighbors that is used in other manifold learning algorithms. Larger datasets usually require a larger perplexity. Consider selecting

s quite

Should be in Macports
py27-scikit-learn @0.15.2 (python, science)

```
def do_tsne_on_vectors( word_vectors ):
    vectors = np.asarray( word_vectors, dtype='float' )

    tsne = TSNE(n_components=2, random_state=42)
    return tsne.fit_transform( vectors )
```



```
def render_json( vectors, labels ):
    output_for_json = []

    vectors = np.array( vectors )

    for i in range( len( vectors ) ):
        new_hash = {}

        new_hash["title"] = str(labels[ i ])
        new_hash["x"] = int(vectors[ i ][ 0 ])
        new_hash["y"] = int(vectors[ i ][ 1 ])

        output_for_json.append( new_hash )

    print json.dumps( output_for_json )
```

Last login: Sun Mar 22 09:53:41 on console
n145-p213:~ hendrikheuer\$ cd Desktop/word2vec_demo/
n145-p213:word2vec_demo hendrikheuer\$ ls
html kafka_metamorphosis.txt
kafka_before_the_law.txt python
n145-p213:word2vec_demo hendrikheuer\$ cd python/
n145-p213:python hendrikheuer\$ ls
data main.py
n145-p213:python hendrikheuer\$ python main.py
[{"y": 2843, "x": -2503, "title": "franz"}, {"y": -298, "x": -2276, "title": "kafka"}, {"y": 2296, "x": 3441, "title": "before"}, {"y": 2354, "x": 2718, "title": "the"}, {"y": 1652, "x": -312, "title": "law"}, {"y": -1005, "x": -920, "title": "sits"}, {"y": -3054, "x": -1026, "title": "a"}, {"y": 3522, "x": 63, "title": "gatekeeper"}, {"y": -268, "x": 548, "title": "to"}, {"y": 70, "x": 2267, "title": "this"}, {"y": 85, "x": 8, "title": "comes"}, {"y": -936, "x": 3024, "title": "man"}, {"y": 1742, "x": -594, "title": "from"}, {"y": -2284, "x": -1748, "title": "country"}, {"y": 2200, "x": 1794, "title": "who"}, {"y": -1851, "x": 1055, "title": "asks"}, {"y": -3701, "x": 831, "title": "gain"}, {"y": 1007, "x": 1876, "title": "entry"}, {"y": -2424, "x": 1749, "title": "into"}, {"y": -2808, "x": 794, "title": "but"}, {"y": 2327, "x": -4408, "title": "says"}, {"y": 1049, "x": -121, "title": "that"}, {"y": 1504, "x": 790, "title": "he"}, {"y": 1122, "x": 47, "title": "can"}, {"y": -1315, "x": 2109, "title": "not"}, {"y": 644, "x": 1454, "title": "grant"}, {"y": 626, "x": 1307, "title": "him"}, {"y": -744, "x": 1846, "title": "at"}, {"y": -1200, "x": -1103, "title": "moment"}, {"y": 164, "x": -3534, "title": "thinks"}, {"y": -2132, "x": 1702, "title": "about"}, {"y": -1773, "x": -811, "title": "it"}, {"y": -650, "x": 2471, "title": "and"}, {"y": 2759, "x": 414, "title": "then"}, {"y": -325, "x": 2099, "title": "if"}, {"y": 3451, "x": 226, "title": "will"}, {"y": -1399, "x": 344, "title": "be"}, {"y": -2626, "x": -1122, "title": "allowed"}, {"y": -2158, "x": 925, "title": "come"}, {"y": -1321, "x": 2718, "title": "in"}, {"y": -776, "x": 1422, "title": "sometime"}, {"y": -1411, "x": 1603, "title": "later"}, {"y": -552, "x": 1620, "title": "on"}, {"y": -1827, "x": 2222, "title": "is"}, {"y": 3054, "x": 2766, "title": "possible"}, {"y": -2603, "x": -329, "title": "now"}, {"y": -659, "x": 280, "title": "gate"}, {"y": -147, "x": -898, "title": "stands"}, {"y": -2500, "x": 1308, "title": "open"}, {"y": 340, "x": 679, "title": "as"}, {"y": 896, "x": 2276, "title": "always"}, {"y": 1451, "x": -646, "title": "walks"}, {"y": -2868, "x": -369, "title": "side"}, {"y": 3440, "x": 1367, "title": "so"}, {"y": 794, "x": 606, "title": "bends"}, {"y": -2368, "x": -90, "title": "over"}, {"y": -270, "x": 1807, "title": "order"}, {"y": 2914, "x": -33, "title": "see"}, {"y": -1854, "x": 1427, "title": "through"}, {"y": -2013, "x": 2158, "title": "inside"}, {"y": -1247, "x": -640, "title": "whe"}, {"y": 196, "x": -2492, "title": "notices"}, {"y": 80, "x": -1434, "title": "laughs"}, {"y": 243, "x": -1867, "title": "you"}, {"y": 2672, "x": 161, "title": "much"}, {"y": -466, "x": 331, "title": "try"}, {"y": 56, "x": 711, "title": "going"}, {"y": 1846, "x": 2343, "title": "spite"}, {"y": 1654, "x": 876, "title": "of"}, {"y": -665, "x": -1047, "title": "my"}, {"y": 359, "x": 1741, "title": "prohibition"}, {"y": -2933, "x": 1245, "title": "take"}, {"y": 2003, "x": -173, "title": "note"}, {"y": -2517, "x": 666, "title": "i"}, {"y": 2847, "x": 434, "title": "am"}, {"y": -62, "x": -2787, "title": "powerful"}, {"y": -2179, "x": 635, "title": "only"}, {"y": 3185, "x": 1373, "title": "most"}, {"y": 914, "x": -2616, "title": "lowly"}, {"y": -901, "x": -2994, "title": "room"}, {"y": }]

D3.js - Data-Driven Docur d3js.org

Overview Examples Documentation Source

Fork me on GitHub

D3 Data-Driven Documents

D3.js is a JavaScript library for manipulating documents based on data. **D3** helps you bring data to life using HTML, SVG, and CSS. D3's emphasis on web standards gives you the full capabilities of modern browsers without tying yourself to a proprietary framework, combining powerful visualization components and a data-driven approach to DOM manipulation.

See more examples.

Target Word	BoW5	BoW2	DEPS
batman	nightwing aquaman catwoman superman manhunter	superman superboy aquaman catwoman batgirl	superman superboy supergirl catwoman aquaman
hogwarts	dumbledore hallows half-blood malfoy snape	evernight sunnydale garderobe blandings collinwood	sunnydale collinwood calarts greendale millfield
turing	nondeterministic non-deterministic computability deterministic finite-state	non-deterministic finite-state nondeterministic buchi primality	pauling hotelling heting lessing hamming
florida	gainesville fla jacksonville tampa lauderdale	fla alabama gainesville tallahassee texas	texas louisiana georgia california carolina
object-oriented	aspect-oriented smalltalk event-driven prolog domain-specific	aspect-oriented event-driven objective-c dataflow 4gl	event-driven domain-specific rule-based data-driven human-centered
dancing	singing dance dances dancers tap-dancing	singing dance dances breakdancing clowning	singing rapping breakdancing miming busking

Discussion:

Can anybody here think of ways
this might help her or him?

word2vec
From theory to practice

Hendrik Heuer
Stockholm NLP Meetup

Further Reading

- Tomas Mikolov, Kai Chen, Greg Corrado, and Jeffrey Dean. Efficient Estimation of Word Representations in Vector Space. In Proceedings of Workshop at ICLR, 2013.
- Tomas Mikolov, Ilya Sutskever, Kai Chen, Greg Corrado, and Jeffrey Dean. Distributed Representations of Words and Phrases and their Compositionality. In Proceedings of NIPS, 2013.
- Tomas Mikolov, Wen-tau Yih, and Geoffrey Zweig. Linguistic Regularities in Continuous Space Word Representations. In Proceedings of NAACL HLT, 2013.

Further Reading

- Richard Socher - Deep Learning for NLP (without Magic)
http://lxmels.it.pt/2014/?page_id=5
- Wang - Introduction to Word2vec and its application to find predominant word senses
<http://compling.hss.ntu.edu.sg/courses/hg7017/pdf/word2vec%20and%20its%20application%20to%20wsd.pdf>
- Exploiting Similarities among Languages for Machine Translation, Tomas Mikolov, Quoc V. Le, Ilya Sutskever,
<http://arxiv.org/abs/1309.4168>
- Title Image by Hans Arp

Further Coding

- word2vec
<https://code.google.com/p/word2vec/>
- word2vec for MacOSX Mavericks
<https://github.com/h10r/word2vec-macosx-maverics>
- Gensim Python Library
<https://radimrehurek.com/gensim/index.html>
- Gensim Tutorials
<https://radimrehurek.com/gensim/tutorial.html>
- Scikit-Learn TSNE
<http://scikit-learn.org/stable/modules/generated/sklearn.manifold.TSNE.html>



word2vec
From theory to practice

Hendrik Heuer
Stockholm NLP Meetup